

Name - Jinal Vyas ASU ID - 1233053785

## Model Training and Evaluation for Segmentation of Glacier Images

Here, I am keeping the initial code from part-2, because it includes

- ✓ comparing three models, and i have done in-detailed post analysis to find the best of the three performing models.

For best visualization of geomaps, i recommend viewing this file in google colab because of its support for google earth engine, however i have plotted the final output files and required plot using matplotlib, so it is platform independent only.

So here, for segmentation, I am using UNet as a base model with three different encoders: ResNet34, ResNet50 and EfficientNet-b4 and Dice Loss as the Evaluation Metric.

Explanation of model choice: ResNet34, ResNet50, and EfficientNet-B4 were chosen as they are popular backbone architectures with varying complexities, enabling a comparison of performance based on model depth and efficiency:

ResNet34: A relatively lightweight model suitable for smaller datasets with faster training and inference. ResNet50: A deeper model with higher capacity, useful for more complex datasets but prone to overfitting. EfficientNet-B4: A highly efficient model that balances accuracy and computational cost using compound scaling.

Dice Loss was chosen as the evaluation metric because: It directly measures the overlap between the predicted and ground truth segmentation masks. It is particularly effective for imbalanced datasets, ensuring small regions are not overlooked. Dice Loss focuses on pixel-level accuracy, making it ideal for image segmentation tasks where precise boundaries are critical. By optimizing Dice Loss, the model maximizes overlap and minimizes false positives/negatives, ensuring better segmentation performance.

```
1 # Install the Rasterio library, which is used for reading and writing geospatial raster data
2 !pip install rasterio

Collecting rasterio
  Downloading rasterio-1.4.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (9.1 kB)
Collecting affine (from rasterio)
  Downloading affine-2.4.0-py3-none-any.whl.metadata (4.0 kB)
Requirement already satisfied: attrs in /usr/local/lib/python3.10/dist-packages (from rasterio) (24.2.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (from rasterio) (2024.8.24)
Requirement already satisfied: click>=4.0 in /usr/local/lib/python3.10/dist-packages (from rasterio) (8.1.7)
Collecting cligj>=0.5 (from rasterio)
  Downloading cligj-0.7.2-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: numpy>=1.24 in /usr/local/lib/python3.10/dist-packages (from rasterio) (1.26.5)
Collecting click-plugins (from rasterio)
  Downloading click_plugins-1.1.1-py2.py3-none-any.whl.metadata (6.4 kB)
Requirement already satisfied: pyparsing in /usr/local/lib/python3.10/dist-packages (from rasterio) (3.2.0)
  Downloading rasterio-1.4.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (22.2 MB)
    22.2/22.2 MB 25.1 MB/s eta 0:00:00
  Downloading cligj-0.7.2-py3-none-any.whl (7.1 kB)
  Downloading affine-2.4.0-py3-none-any.whl (15 kB)
  Downloading click_plugins-1.1.1-py2.py3-none-any.whl (7.5 kB)
  Installing collected packages: cligj, click-plugins, affine, rasterio
  Successfully installed affine-2.4.0 click-plugins-1.1.1 cligj-0.7.2 rasterio-1.4.3
```

```
1 # Install the segmentation_models_pytorch library for pretrained segmentation model, U-Net in this cell
2 !pip install segmentation_models_pytorch

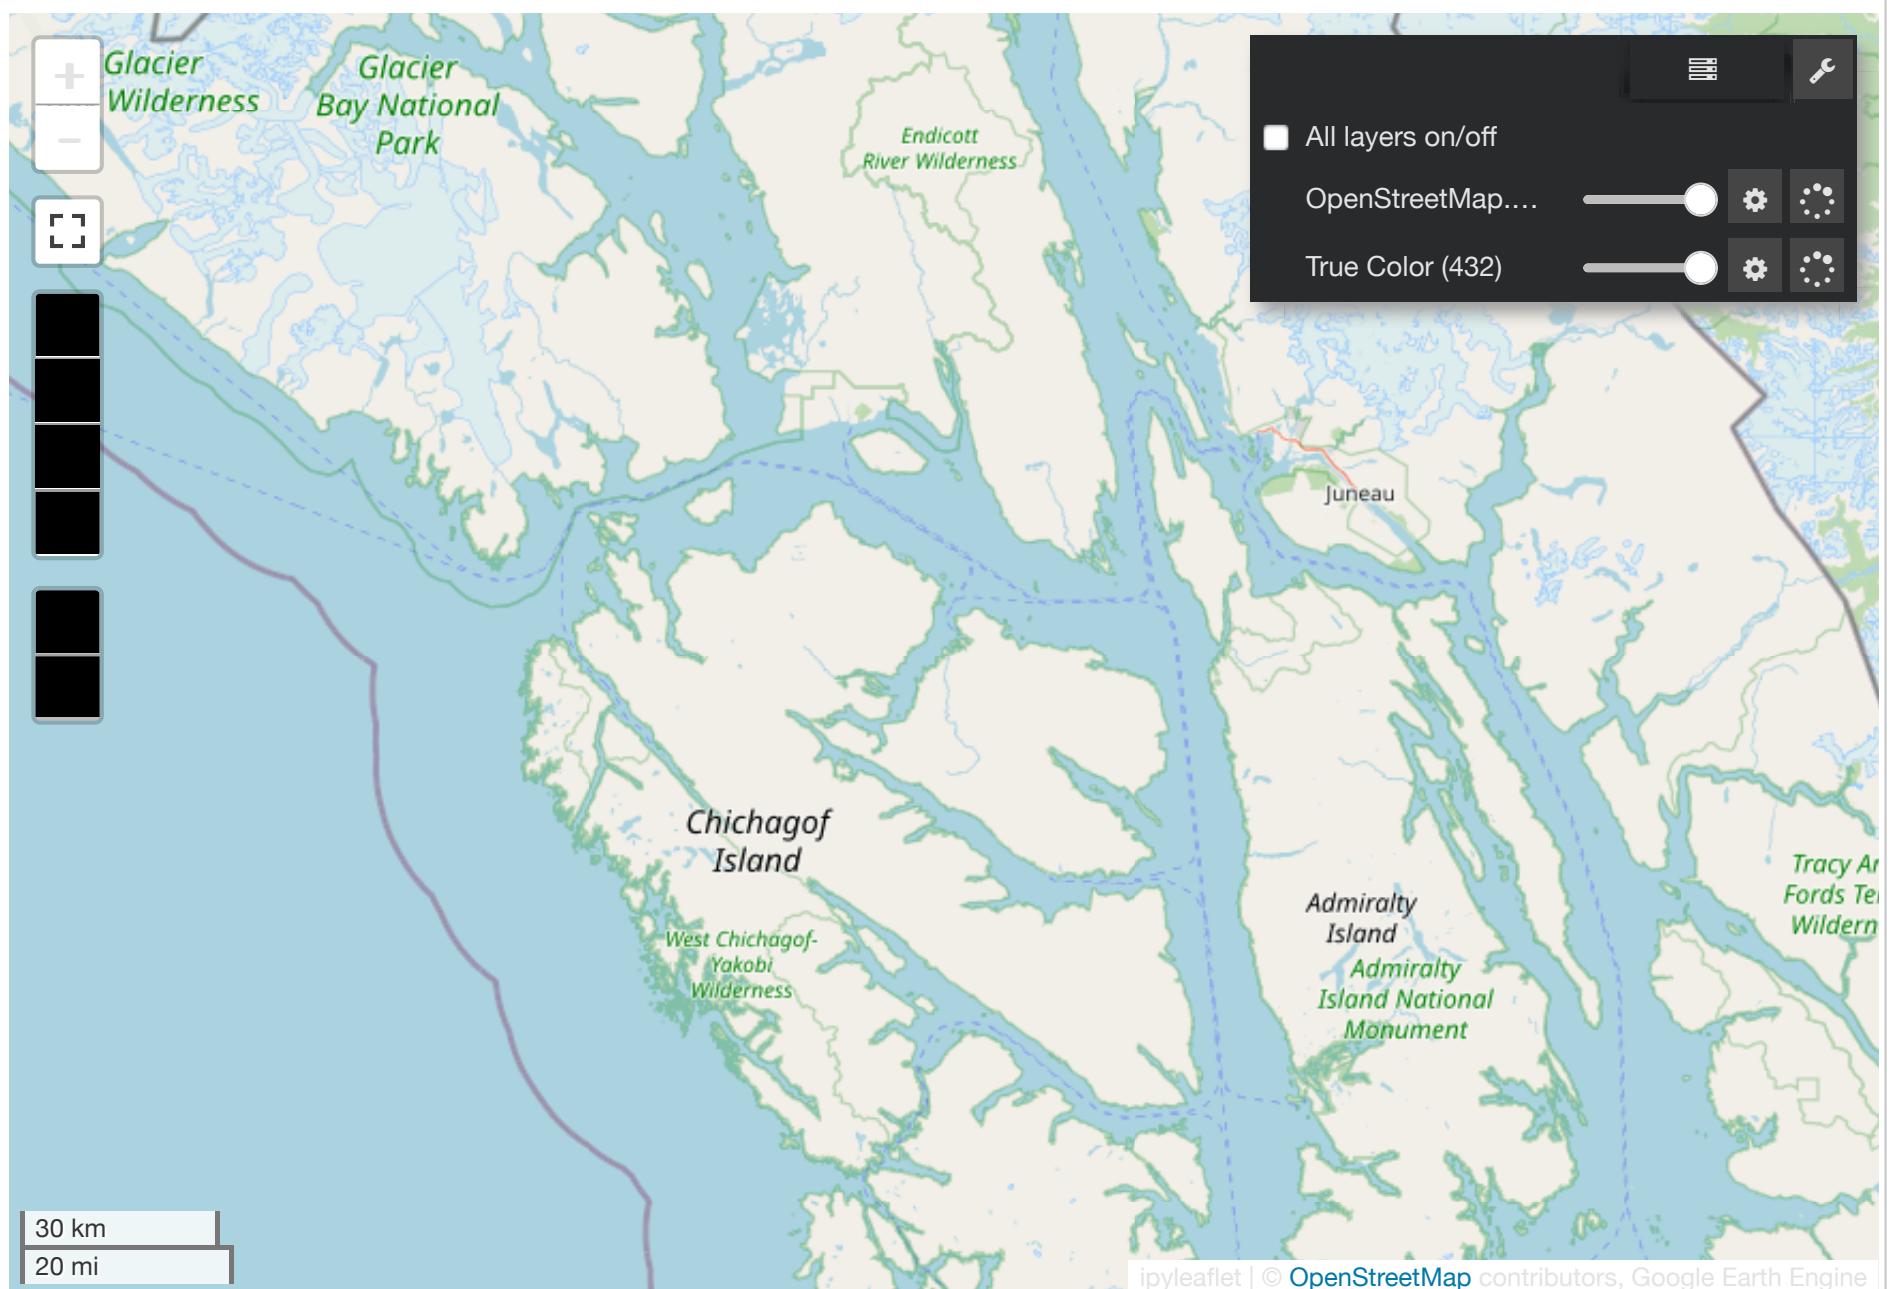
Collecting segmentation_models_pytorch
  Downloading segmentation_models_pytorch-0.3.4-py3-none-any.whl.metadata (30 kB)
Collecting efficientnet-pytorch==0.7.1 (from segmentation_models_pytorch)
  Downloading efficientnet_pytorch-0.7.1.tar.gz (21 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: huggingface-hub>=0.24.6 in /usr/local/lib/python3.10/dist-packages (from segmentation_models_pytorch)
Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from segmentation_models_pytorch)
Collecting pretrainedmodels==0.7.4 (from segmentation_models_pytorch)
  Downloading pretrainedmodels-0.7.4.tar.gz (58 kB)
  58.8/58.8 KB 4.4 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from segmentation_models_pytorch)
Collecting timm==0.9.7 (from segmentation_models_pytorch)
  Downloading timm-0.9.7-py3-none-any.whl.metadata (58 kB)
  58.8/58.8 KB 5.0 MB/s eta 0:00:00
Requirement already satisfied: torchvision>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from segmentation_models_pytorch)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from segmentation_models_pytorch)
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from efficientnet-pytorch)
Collecting munch (from pretrainedmodels==0.7.4->segmentation_models_pytorch)
  Downloading munch-4.0.0-py2.py3-none-any.whl.metadata (5.9 kB)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from timm==0.9.7->segmentation_models_pytorch)
Requirement already satisfied: safetensors in /usr/local/lib/python3.10/dist-packages (from timm==0.9.7->segmentation_models_pytorch)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from huggingface-hub)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from timm==0.9.7->segmentation_models_pytorch)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torchvision>=0.5.0->segmentation_models_pytorch)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch->efficientnet-pytorch)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch->efficientnet-pytorch)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch->efficientnet-pytorch)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->huggingface-hub)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->segmentation_models_pytorch)
  Downloading segmentation_models_pytorch-0.3.4-py3-none-any.whl (109 kB)
  109.5/109.5 KB 9.6 MB/s eta 0:00:00
  Downloading timm-0.9.7-py3-none-any.whl (2.2 MB)
  2.2/2.2 MB 31.6 MB/s eta 0:00:00
  Downloading munch-4.0.0-py2.py3-none-any.whl (9.9 kB)
Building wheels for collected packages: efficientnet-pytorch, pretrainedmodels
  Building wheel for efficientnet-pytorch (setup.py) ... done
  Created wheel for efficientnet-pytorch: filename=efficientnet_pytorch-0.7.1-py3-none-any.whl size=1642499
  Stored in directory: /root/.cache/pip/wheels/03/3f/e9/911b1bc46869644912bda90a56bcf7b960f20b5187feea3ba
  Building wheel for pretrainedmodels (setup.py) ... done
  Created wheel for pretrainedmodels: filename=pretrainedmodels-0.7.4-py3-none-any.whl size=60944 sha256=5a2...
  Stored in directory: /root/.cache/pip/wheels/35/cb/a5/8f534c60142835bfc889f9a482e4a67e0b817032d9c6883b0
Successfully built efficientnet-pytorch pretrainedmodels
Installing collected packages: munch, efficientnet-pytorch, timm, pretrainedmodels, segmentation_models_pytorch
  Attempting uninstall: timm
    Found existing installation: timm 1.0.12
    Uninstalling timm-1.0.12:
      Successfully uninstalled timm-1.0.12
Successfully installed efficientnet-pytorch-0.7.1 munch-4.0.0 pretrainedmodels-0.7.4 segmentation_models_pytorch-0.3.4
```

```
1 # installing required libraries
2 import ee
3 import geemap.core as geemap
4 import torch
5 import rasterio
6 from torch.utils.data import Dataset, DataLoader
7 from torchvision import transforms
8 import numpy as np
9 import matplotlib.pyplot as plt
10 import numpy as np
11 from osgeo import gdal
12 import seaborn as sns
13 import time
14 from tqdm import tqdm
15 import cv2
16 from google.colab.patches import cv2_imshow
17 import os
18 from PIL import Image
19 import segmentation_models_pytorch as smp
20 from sklearn.model_selection import KFold
```

```
1 # authenticating the access of google earth engine
2 ee.Authenticate()
3 ee.Initialize(project='ee-jinalvyasict19')
```

```
1 # Landsat-8 was launched in february, 2013, hence am considering data of past 10 years starting from
2 # I am combining the data of every month into a single image through its median, hence i am generating
3 # However, due to cloud mask, 38 images were found completely empty/corrupted, so i dropped them and
4 # 10% of images among those 82 may have very less data but still I am using it for training and testing
5 dataset = ee.ImageCollection('LANDSAT/LC08/C02/T1_L2').filterDate(
6     '2023-06-01', '2023-12-01'
7 )
8 # masking alaska's mendenhall glacier coordinates
9 alaska_bounds = ee.Geometry.Polygon([[-134.62, 58.42], [-134.62, 58.59], [-134.42, 58.59], [-134.42, 58.42],
10
11 # selecting the appropriate bands required, in the data preprocessing part, i have indetailed explanation
12 # Optical Bands (SR_B3 to SR_B7): Represent visible, near-infrared (NIR), and shortwave infrared (SWIR)
13 # Thermal Bands (ST_B.*): Capture surface temperature, useful for monitoring heat and environmental
14 def apply_scale_factors(image):
15     global alaska_bounds
16     optical_bands = image.select(['SR_B3', 'SR_B4', 'SR_B5', 'SR_B6', 'SR_B7']).multiply(0.0000275).add(-273.15)
17     thermal_bands = image.select('ST_B.*').multiply(0.00341802).add(149.0)
18     image_with_bands = image.addBands(optical_bands, None, True).addBands(
19         thermal_bands, None, True
20     )
21     return image_with_bands.clip(alaska_bounds)
22
23 # clouds many times distort the true images, hence it is crucial to apply cloud mask and remove the
24 def mask_clouds_shadows(image):
25     qa_mask = image.select('QA_PIXEL')
26     cloud_mask = (qa_mask.bitwiseAnd(1 << 3)).eq(0)
27     shadow_mask = (qa_mask.bitwiseAnd(1 << 4)).eq(0)
28     mask = cloud_mask.And(shadow_mask)
29     return image.updateMask(mask)
30
31 dataset = dataset.map(apply_scale_factors)
32 dataset = dataset.map(mask_clouds_shadows)
33
34 visualization = {
35     'bands': ['SR_B4', 'SR_B3', 'SR_B2'],
36     'min': 0.0,
37     'max': 0.3,
38 }
39
40 map = geemap.Map()
41 map.centerObject(alaska_bounds, 10)
42 map.add_layer(dataset, visualization, 'True Color (432)')
```

43 map



```

1 # applying preprocessing to the images of past ten years i.e. ranging from January, 2014 to December
2 images = {}
3
4 for year in tqdm(range(2014, 2024)):
5     for month in tqdm(range(1, 13)):
6         if month != 2:
7             dataset = ee.ImageCollection('LANDSAT/LC08/C02/T1_L2').filterDate(
8                 rf'{year}-{month}-01', rf'{year}-{month}-30'
9             )
10        else:
11            dataset = ee.ImageCollection('LANDSAT/LC08/C02/T1_L2').filterDate(
12                rf'{year}-{month}-01', rf'{year}-{month}-28'
13            )
14        dataset = dataset.map(apply_scale_factors)
15        dataset = dataset.map(mask_clouds_shadows)
16        # the reason I am taking median and not all the images of every month is, the frequency of Landsat
17        # and also there is a good probability that either of those image may be corrupted due to clouds
18        # as median will also be less affected by the outliers (which may be the NaN values recorded due
19        images[f'{year}-{month}'] = dataset.median()
20 images

```

0% | 0/10 [00:00<?, ?it/s]  
100% | ██████████ | 12/12 [00:00<00:00, 216.29it/s]

100% | ██████████ | 12/12 [00:00<00:00, 257.08it/s]  
20% | █ | 2/10 [00:00<00:00, 16.57it/s]  
100% | ██████████ | 12/12 [00:00<00:00, 129.53it/s]

0% | 0/12 [00:00<?, ?it/s]  
100% | ██████████ | 12/12 [00:00<00:00, 99.21it/s]  
40% | █ | 4/10 [00:00<00:00, 10.70it/s]  
0% | 0/12 [00:00<?, ?it/s]  
100% | ██████████ | 12/12 [00:00<00:00, 55.89it/s]

0% | 0/12 [00:00<?, ?it/s]  
100% | ██████████ | 12/12 [00:00<00:00, 74.43it/s]  
60% | █ | 6/10 [00:00<00:00, 6.84it/s]  
0% | 0/12 [00:00<?, ?it/s]

```
100%|██████████| 12/12 [00:00<00:00, 55.33it/s]
70%|██████| 7/10 [00:01<00:00, 5.86it/s]
0%|          | 0/12 [00:00<?, ?it/s]
100%|██████████| 12/12 [00:00<00:00, 86.09it/s]
80%|██████| 8/10 [00:01<00:00, 5.82it/s]
100%|██████████| 12/12 [00:00<00:00, 125.52it/s]
90%|███████| 9/10 [00:01<00:00, 6.44it/s]
0%|          | 0/12 [00:00<?, ?it/s]
100%|██████████| 12/12 [00:00<00:00, 53.14it/s]
100%|██████████| 10/10 [00:01<00:00, 6.48it/s]
{'2014-1': <ee.image.Image at 0x7d7b9dc72fe0>,
 '2014-2': <ee.image.Image at 0x7d7b9d70c040>,
 '2014-3': <ee.image.Image at 0x7d7b9dcab130>,
 '2014-4': <ee.image.Image at 0x7d7b9dcaab60>,
 '2014-5': <ee.image.Image at 0x7d7b9dd15c00>,
 '2014-6': <ee.image.Image at 0x7d7b9dd4b070>,
 '2014-7': <ee.image.Image at 0x7d7b9d2352d0>,
 '2014-8': <ee.image.Image at 0x7d7b9d236800>,
 '2014-9': <ee.image.Image at 0x7d7b9d2342b0>,
 '2014-10': <ee.image.Image at 0x7d7b9d29c640>,
 '2014-11': <ee.image.Image at 0x7d7b9d29d6f0>,
 '2014-12': <ee.image.Image at 0x7d7b9d29e6e0>,
 '2015-1': <ee.image.Image at 0x7d7b9dd04940>,
 '2015-2': <ee.image.Image at 0x7d7b9dca9a50>,
 '2015-3': <ee.image.Image at 0x7d7b9dd061d0>,
 '2015-4': <ee.image.Image at 0x7d7b9dc62770>,
 '2015-5': <ee.image.Image at 0x7d7b9dc62260>,
 '2015-6': <ee.image.Image at 0x7d7b9dc63e20>,
 '2015-7': <ee.image.Image at 0x7d7b9de7ea40>,
 '2015-8': <ee.image.Image at 0x7d7b9d724640>,
 '2015-9': <ee.image.Image at 0x7d7b9d726e60>,
 '2015-10': <ee.image.Image at 0x7d7b9d726140>,
 '2015-11': <ee.image.Image at 0x7d7b9d724b20>,
 '2015-12': <ee.image.Image at 0x7d7b9c49c520>,
 '2016-1': <ee.image.Image at 0x7d7b9c49d600>,
 '2016-2': <ee.image.Image at 0x7d7b9c49e5f0>,
 '2016-3': <ee.image.Image at 0x7d7b9c49f310>,
 '2016-4': <ee.image.Image at 0x7d7b9c4ac220>,
 '2016-5': <ee.image.Image at 0x7d7b9c4ad210>,
 '2016-6': <ee.image.Image at 0x7d7b9c4ae200>,
 '2016-7': <ee.image.Image at 0x7d7b9c4af130>,
 '2016-8': <ee.image.Image at 0x7d7b9c4afee0>,
 '2016-9': <ee.image.Image at 0x7d7b9c4b8f10>,
 '2016-10': <ee.image.Image at 0x7d7b9c4b9f00>,
 '2016-11': <ee.image.Image at 0x7d7b9c4babf0>,
 '2016-12': <ee.image.Image at 0x7d7b9c4bbbb0>,
 '2017-1': <ee.image.Image at 0x7d7b9c4c8cd0>,
 '2017-2': <ee.image.Image at 0x7d7b9c4c9a80>,
 '2017-3': <ee.image.Image at 0x7d7b9c4cab0>,
 '2017-4': <ee.image.Image at 0x7d7b9c4cb8b0>,
 '2017-5': <ee.image.Image at 0x7d7b9c4d8850>,
 '2017-6': <ee.image.Image at 0x7d7b9c4d9840>,
 '2017-7': <ee.image.Image at 0x7d7b9c4da830>,
 '2017-8': <ee.image.Image at 0x7d7b9c4db820>,
 '2017-9': <ee.image.Image at 0x7d7b9c4e41f0>,
 '2017-10': <ee.image.Image at 0x7d7b9c4e5480>,
 '2017-11': <ee.image.Image at 0x7d7b9c4e6470>,
 '2017-12': <ee.image.Image at 0x7d7b9c4e7460>,
 '2018-1': <ee.image.Image at 0x7d7b9c4f4040>,
 '2018-2': <ee.image.Image at 0x7d7b9c4f5360>,
 '2018-3': <ee.image.Image at 0x7d7b9c4f6320>,
 '2018-4': <ee.image.Image at 0x7d7b9c4f7310>,
 '2018-5': <ee.image.Image at 0x7d7b9c504340>,
 '2018-6': <ee.image.Image at 0x7d7b9c504c10>,
 '2018-7': <ee.image.Image at 0x7d7b9c505f60>,
 '2018-8': <ee.image.Image at 0x7d7b9c506f50>,
 '2018-9': <ee.image.Image at 0x7d7b9c507f40>,
 '2018-10': <ee.image.Image at 0x7d7b9c518bb0>,
 '2018-11': <ee.image.Image at 0x7d7b9c519c60>,
 '2018-12': <ee.image.Image at 0x7d7b9c51ac50>,
 '2019-1': <ee.image.Image at 0x7d7b9c51ae90>,
 '2019-2': <ee.image.Image at 0x7d7b9c528d60>,
 '2019-3': <ee.image.Image at 0x7d7b9c529990>,
 '2019-4': <ee.image.Image at 0x7d7b9c52a980>,
 '2019-5': <ee.image.Image at 0x7d7b9c52b970>,
 '2019-6': <ee.image.Image at 0x7d7b9c5349a0>,
 '2019-7': <ee.image.Image at 0x7d7b9c535990>,
 '2019-8': <ee.image.Image at 0x7d7b9c5366b0>,
 '2019-9': <ee.image.Image at 0x7d7b9c537790>,
```

```
'2019-10': <ee.image.Image at 0x7d7b9c54c640>,
'2019-11': <ee.image.Image at 0x7d7b9c54d630>,
'2019-12': <ee.image.Image at 0x7d7b9c54e410>,
'2020-1': <ee.image.Image at 0x7d7b9c554220>,
'2020-2': <ee.image.Image at 0x7d7b9c54f7f0>,
'2020-3': <ee.image.Image at 0x7d7b9c5553c0>,
'2020-4': <ee.image.Image at 0x7d7b9c5560b0>,
'2020-5': <ee.image.Image at 0x7d7b9c557130>,
'2020-6': <ee.image.Image at 0x7d7b9c568160>,
'2020-7': <ee.image.Image at 0x7d7b9c5691b0>,
'2020-8': <ee.image.Image at 0x7d7b9c56a4d0>,
'2020-9': <ee.image.Image at 0x7d7b9c56b1f0>,
'2020-10': <ee.image.Image at 0x7d7b9c5740d0>,
'2020-11': <ee.image.Image at 0x7d7b9c5750f0>,
'2020-12': <ee.image.Image at 0x7d7b9c5760e0>,
'2021-1': <ee.image.Image at 0x7d7b9c576dd0>,
'2021-2': <ee.image.Image at 0x7d7b9c577dc0>,
'2021-3': <ee.image.Image at 0x7d7b9c38cdf0>,
'2021-4': <ee.image.Image at 0x7d7b9c38dde0>,
'2021-5': <ee.image.Image at 0x7d7b9c38edd0>,
'2021-6': <ee.image.Image at 0x7d7b9c38f550>,
'2021-7': <ee.image.Image at 0x7d7b9c398a60>,
'2021-8': <ee.image.Image at 0x7d7b9c399a50>,
'2021-9': <ee.image.Image at 0x7d7b9c39aa40>,
'2021-10': <ee.image.Image at 0x7d7b9c39ba30>,
'2021-11': <ee.image.Image at 0x7d7b9c3a8760>,
'2021-12': <ee.image.Image at 0x7d7b9c3a9750>,
'2022-1': <ee.image.Image at 0x7d7b9c3aa890>,
'2022-2': <ee.image.Image at 0x7d7b9c3ab6d0>,
'2022-3': <ee.image.Image at 0x7d7b9c3bc610>,
'2022-4': <ee.image.Image at 0x7d7b9c3bd600>,
'2022-5': <ee.image.Image at 0x7d7b9c3be5f0>,
'2022-6': <ee.image.Image at 0x7d7b9c3bf640>,
'2022-7': <ee.image.Image at 0x7d7b9c3c4250>,
'2022-8': <ee.image.Image at 0x7d7b9c3c5300>,
'2022-9': <ee.image.Image at 0x7d7b9c3c62f0>,
'2022-10': <ee.image.Image at 0x7d7b9c3c72e0>,
'2022-11': <ee.image.Image at 0x7d7b9c3bffa0>,
'2022-12': <ee.image.Image at 0x7d7b9c3dd000>,
'2023-1': <ee.image.Image at 0x7d7b9c3de050>,
'2023-2': <ee.image.Image at 0x7d7b9c3df0d0>,
'2023-3': <ee.image.Image at 0x7d7b9c3df970>,
'2023-4': <ee.image.Image at 0x7d7b9c3eceb0>,
'2023-5': <ee.image.Image at 0x7d7b9c3edea0>,
'2023-6': <ee.image.Image at 0x7d7b9c3eee90>,
'2023-7': <ee.image.Image at 0x7d7b9c3eff40>,
'2023-8': <ee.image.Image at 0x7d7b9c3f4970>,
```

```
1 # mounting google drive, for easy storage and access of data
2 from google.colab import drive
3 drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
1 # Creating tasks for downloading the images to a folder named, 'Alasks Dataset', in the drive, it is
2 for date, image in tqdm(images.items()):
3     export_task = ee.batch.Export.image.toDrive(
4         image=image,
5         description=f'image_{date}',
6         folder='Alaska_Dataset',
7         region=alaska_bounds,
8         scale=30,
9         fileFormat = 'GeoTIFF',
10        maxPixels=1e13
11    )
12    export_task.start()
13    while export_task.active():
14        pass
```

100% |██████████| 120/120 [1:16:02<00:00, 38.02s/it]

```
1 # Here, using the tiff files derived from the above task, I am extracting the bands 3 to 7, replacing
2 # Then, using bands 3, 4 and 5, I am plotting RGB images, and also NDSI ( Normalized Difference Snow
3 # Here, as visually representable, RGB images better distinguish the glacier from NDSI images, hence
4
5 # I am also applying thresholding of rgb_image.max() == 1, such that months producing empty tif file
6 for files in images.keys():
7     file_path = f"/content/drive/My Drive/Alaska_Dataset/image_{files}.tif"
8     try:
9         with rasterio.open(file_path) as src:
10             band_3 = src.read(3)
11             band_4 = src.read(4)
12             band_5 = src.read(5)
13             band_6 = src.read(6)
14             band_7 = src.read(7)
15
16             band_3 = np.nan_to_num(band_3, nan=0)
17             band_4 = np.nan_to_num(band_4, nan=0)
18             band_5 = np.nan_to_num(band_5, nan=0)
19             band_6 = np.nan_to_num(band_6, nan=0)
20             band_7 = np.nan_to_num(band_7, nan=0)
21
22             band_3_flat = band_3.flatten()
23             band_4_flat = band_4.flatten()
24             band_5_flat = band_5.flatten()
25             band_6_flat = band_6.flatten()
26             band_7_flat = band_7.flatten()
27             band_5_norm = band_5 / band_5.max()
28             band_4_norm = band_4 / band_4.max()
29             band_3_norm = band_3 / band_3.max()
30
31             rgb_image = np.stack((band_5_norm, band_4_norm, band_3_norm), axis=-1)
32             NDSI = (band_3 - band_6) / (band_3 + band_6)
33             if rgb_image.max() == 1:
34                 plt.figure(figsize=(10, 10))
35                 plt.imshow(rgb_image)
36                 plt.title("False Color Composite (NIR, Red, Green)")
37                 plt.axis('off')
38                 plt.savefig(f"/content/drive/My Drive/Alaska_Dataset/image_{files}.png")
39                 plt.show()
40                 plt.figure(figsize=(7, 7))
41                 plt.imshow(NDSI, cmap='Spectral')
42                 plt.title('Normalized Difference Snow Index (NDSI)')
43                 plt.colorbar()
44                 plt.axis('off')
45                 plt.clim(-5, 1)
46                 plt.savefig(f"/content/drive/My Drive/Alaska_Dataset/image_{files}_NDSI.png")
47                 plt.show()
48     except Exception as e:
49         print(e)
```

```

1 # I am creating binary masks using the RGB images and OpenCV, first I am converting them to grayscale
2 # and the non-glacier region, including land and vegetation will be marked as black. I am also cropping
3 # as it can cause discrepancy in the segmentation.
4 for files in images.keys():
5     file_path = rf"/content/drive/My Drive/Alaska_Dataset/image_{files}.png"
6     if not os.path.exists(file_path):
7         print(f'image_{files}.png was not saved, because it contained no data!')
8         continue
9     file_path = rf"/content/drive/My Drive/Alaska_Dataset/image_{files}.png"
10    img = cv2.imread(file_path, cv2.IMREAD_GRAYSCALE)
11    # print(img)
12    _, mask = cv2.threshold(img, 128, 255, cv2.THRESH_BINARY)
13    # cv2.imwrite(f"/content/drive/My Drive/Alaska_Dataset/mask_{files}.png", mask)
14    x, y, w, h = 55, 82, 400, 340
15    cropped_mask = mask[y:y+h, x:x+w]
16    cv2.imshow(cropped_mask)
17    cv2.waitKey(0)
18    cv2.destroyAllWindows()
19    cv2.imwrite(f"/content/drive/My Drive/Alaska_Dataset/mask_{files}.png", cropped_mask)

```

```

1 # I am creating a custom dataset such that it applies transformation to the images and masks, here I
2 # modified, for example, larger models will function better on 512*512 images, and rather overfit on
3 # hence 224*224 seems like an appropriate choice.
4 class SegmentationDataset(Dataset):
5     def __init__(self, image_path, mask_path, transform=None):
6         self.image_paths = image_path
7         self.mask_paths = mask_path
8         self.transform = transforms.Compose([
9             transforms.Resize((224, 224)), # Resize the image
10            transforms.ToTensor(), # Convert to tensor
11            transforms.Normalize(mean=[0.485, 0.456, 0.406],
12                                std=[0.229, 0.224, 0.225]) # Normalize using ImageNet stats
13        ])
14        self.mask_transform = transforms.Compose([
15            transforms.Resize((224, 224)), # Resize the image
16            transforms.ToTensor()
17        ])
18    def __len__(self):
19        return len(self.image_paths)
20    def __getitem__(self, idx):
21        image = Image.open(self.image_paths[idx]).convert("RGB")
22        x, y, w, h = 55, 82, 400, 340
23        cropped_image = self.transform(image.crop((x, y, x + w, y + h)))
24        mask = self.mask_transform(Image.open(self.mask_paths[idx]))
25        # print(image.shape, mask.shape)
26        return cropped_image, mask

```

```

1 # I am appending all the rgb image paths and mask paths to two variables, which I will later use for
2 image_paths = []
3 mask_paths = []
4 for files in images.keys():
5     image_path = rf"/content/drive/My Drive/Alaska_Dataset/image_{files}.png"
6     if not os.path.exists(image_path):
7         print(f'image_{files}.png was not saved, because it contained no data!')
8         continue
9     mask_path = rf"/content/drive/My Drive/Alaska_Dataset/mask_{files}.png"
10    image_paths.append(image_path)
11    mask_paths.append(mask_path)
12 image_paths, mask_paths

```

```

image_2014-1.png was not saved, because it contained no data!
image_2014-9.png was not saved, because it contained no data!
image_2014-12.png was not saved, because it contained no data!
image_2015-1.png was not saved, because it contained no data!
image_2015-2.png was not saved, because it contained no data!
image_2015-11.png was not saved, because it contained no data!
image_2015-12.png was not saved, because it contained no data!
image_2016-1.png was not saved, because it contained no data!
image_2016-8.png was not saved, because it contained no data!

```





```
'/content/drive/My Drive/Alaska_Dataset/mask_2019-8.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2019-9.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2019-10.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-2.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-3.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-4.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-5.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-7.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-9.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2020-10.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-2.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-3.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-4.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-5.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-6.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-7.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-8.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-9.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2021-10.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-3.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-4.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-5.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-6.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-8.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-9.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2022-11.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2023-2.png',
'/content/drive/My Drive/Alaska_Dataset/mask_2023-3.png',
```

```
1 # as explained earlier, these are total number of images I got after filtering cloud mask, and take
2 len(image_paths)
```

82

```
1 # defining device type, here I am using T4 GPU provided by google colab.
2 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
3 device
device(type='cuda')
```

```
1 # defining the train function for single epoch, which returns the total loss per epoch i.e. the loss
2 def train(model, train_loader, optimizer, loss, device):
3     model.train()
4     epoch_loss = 0
5     for images, masks in train_loader:
6         images = images.to(device)
7         masks = masks.to(device)
8         outputs = model(images)
9         # print(outputs.shape, masks.shape)
10        batch_loss = loss(outputs, masks)
11        optimizer.zero_grad()
12        batch_loss.backward()
13        optimizer.step()
14        epoch_loss += batch_loss.item()
15    return epoch_loss/len(train_loader)
```

```

1 # defining the testing function, which return the total testing loss per epoch, averaged over the batch
2 def test(model, test_loader, loss, device):
3     model.eval()
4     test_loss = 0
5     with torch.no_grad():
6         for images, masks in test_loader:
7             images = images.to(device)
8             masks = masks.to(device)
9             outputs = model(images)
10            batch_loss = loss(outputs, masks)
11            test_loss += batch_loss.item()
12    return test_loss/len(test_loader)

```

```

1 # this function is used to plot all the images in the test loader. there are three images being plotted
2 def test_image_plot(test_loader, model):
3     for images, masks in test_loader:
4         images = images.to(device)
5         masks = masks.to(device)
6         outputs = model(images)
7         pred_masks = (outputs > 0.5).float()
8         for image, mask, pred_mask in zip(images, masks, pred_masks):
9             fig, axs = plt.subplots(1, 3, figsize=(15, 5))
10            axs[0].imshow(image.cpu().numpy().transpose(1, 2, 0)) # Convert tensor to numpy (C, H, W) to (H, W, C)
11            axs[0].set_title("Original Image")
12            axs[1].imshow(mask.cpu().numpy().squeeze(), cmap='gray') # Mask is single channel, squeeze it
13            axs[1].set_title("Ground Truth Mask")
14            axs[1].axis('off')
15            axs[2].imshow(pred_mask.cpu().numpy().squeeze(), cmap='gray') # Predicted mask (squeeze it)
16            axs[2].set_title("Predicted Mask")
17            axs[2].axis('off')
18            plt.show()

```

```

1 # sorting the images and masks as per date and time, to make sure the indexes of both the lists are same
2 image_paths.sort()
3 mask_paths.sort()
4 image_paths[:5], mask_paths[:5]

([ '/content/drive/My Drive/Alaska_Dataset/image_2014-10.png',
  '/content/drive/My Drive/Alaska_Dataset/image_2014-11.png',
  '/content/drive/My Drive/Alaska_Dataset/image_2014-2.png',
  '/content/drive/My Drive/Alaska_Dataset/image_2014-3.png',
  '/content/drive/My Drive/Alaska_Dataset/image_2014-4.png'],
 [ '/content/drive/My Drive/Alaska_Dataset/mask_2014-10.png',
  '/content/drive/My Drive/Alaska_Dataset/mask_2014-11.png',
  '/content/drive/My Drive/Alaska_Dataset/mask_2014-2.png',
  '/content/drive/My Drive/Alaska_Dataset/mask_2014-3.png',
  '/content/drive/My Drive/Alaska_Dataset/mask_2014-4.png'])

```

```

1 # For testing the model, I am using 20% of the data, i.e. I will divide these indices into 80-20 split
2 # However, keeping one image separate to compare segmentation output mask produced the best weights
3
4 indices = list(range(len(image_paths[:81])))
5 indices

[0,
 1,
 2,
 3,
 4,
 5,
 6,
 7,
 8,
 9,
 10,
 11,
 12,
 13,
 14,
 15,
 16,
 17,
 18,
 19,
 20,
 21,
 22,
 23,
 24,
 25,
 26,
 27,
 28,
 29,
 30,
 31,
 32,
 33,
 34,
 35,
 36,
 37,
 38,
 39,
 40,
 41,
 42,
 43,
 44,
 45,
 46,
 47,
 48,
 49,
 50,
 51,
 52,
 53,
 54,
 55,
 56,
 57,
 58,
 59,
 60,
 61,
 62,
 63,
 64,
 65,
 66,
 67,
 68,
 69,
 70,
 71,
 72,
 73,
 74,
 75,
 76,
 77,
 78,
 79,
 80]

```

```
14,  
15,  
16,  
17,  
18,  
19,  
20,  
21,  
22,  
23,  
24,  
25,  
26,  
27,  
28,  
29,  
30,  
31,  
32,  
33,  
34,  
35,  
36,  
37,  
38,  
39,  
40,  
41,  
42,  
43,  
44,  
45,  
46,  
47,  
48,  
49,  
50,  
51,  
52,  
53,  
54,  
55,  
56,  
57,  
58,  
59,  
60,  
61,  
62,  
63,  
64,  
65,  
66,  
67,  
68,  
69,  
70,  
71,  
72,  
73,  
74,  
75,  
76,  
77,  
78,
```

```
1 # Initialize K-Fold Cross-Validation with 5 splits  
2 # shuffle=True ensures the data is randomly shuffled before splitting  
3 # random_state=42 ensures reproducibility of the splits  
4 k_folds = 5  
5 kf = KFold(n_splits=k_folds, shuffle=True, random_state=42)  
6 kf
```

```
KFold(n_splits=5, random_state=42, shuffle=True)
```

## UNet with Encoder ResNet34 - Training and Evaluation

```

1 # here i will be iterating over 5 folds and in each fold, i am dividing the data into 80-20 split i.
2 # Hence as visible, there are 65-66 training images every iteration and 16-17 test images.
3 # I am globally defining one variable called model_best_loss, which will track when the TEST loss fu
4 # the best model obtained.
5 # I am using resnet34 as my encoder here, with pretrained weights.
6 k_folds_results = []
7 resnet34_best_test_loss = 10000
8 # As explained earlier in the initial cells, I am using Dice Loss as my evaluation metric as it min:
9 loss = smp.losses.DiceLoss(mode='binary')
10 for fold, (train_index, test_index) in enumerate(kf.split(indices)):
11     train_image_paths = [image_paths[i] for i in train_index]
12     train_mask_paths = [mask_paths[i] for i in train_index]
13     test_image_paths = [image_paths[i] for i in test_index]
14     test_mask_paths = [mask_paths[i] for i in test_index]
15     print(len(train_index), len(test_index))
16     train_dataset = SegmentationDataset(train_image_paths, train_mask_paths)
17     test_dataset = SegmentationDataset(test_image_paths, test_mask_paths)
18     train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
19     test_loader = DataLoader(test_dataset, batch_size=8, shuffle=True)
20     model = smp.Unet(
21         encoder_name="resnet34", #encoder defined
22         encoder_weights="imagenet", #imagenet pretrained weights are used
23         in_channels=3, #3 classes as input images are RGB
24         classes=1 #as output is supposed to be a binary mask, hence 1 output classes only.
25     )
26     model = model.to(device)
27     # here I am using the classic Adam as an optimizer with learning rate 1e-4, decreasing it from hei
28     # Adam optimizer was chosen because:
29     # - It combines the benefits of Momentum and RMSprop, adapting the learning rate for each parameter
30     # - It is computationally efficient and works well for large datasets and high-dimensional problems
31     # - Adam performs well with minimal hyperparameter tuning, making it a reliable choice for most tasks
32     # - Its adaptive nature ensures faster convergence and robustness to noisy gradient
33     # And the model didn't overfit with the given hyperparameters, hence i didn't provide weight decay
34     optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)
35     for epoch in tqdm(range(30)):
36         train_loss = train(model, train_loader, optimizer, loss, device)
37         test_loss = test(model, test_loader, loss, device)
38         if test_loss < resnet34_best_test_loss:
39             resnet34_best_test_loss = test_loss
40             torch.save(model.state_dict(), f"Best_model_resnet34.pth")
41             print("\nBest Model Saved with Test Loss: ", resnet34_best_test_loss)
42             best_test_loader = test_loader
43     k_folds_results.append(test_loss)
44 print("\nCross Validation Results: ", k_folds_results)
45 print("\nAverage loss: ", sum(k_folds_results)/len(k_folds_results))

64 17
Downloading: "https://download.pytorch.org/models/resnet34-333f7ec4.pth" to /root/.cache/torch/hub/checkpoints/resnet34-333f7ec4.pth
100%|██████████| 83.3M/83.3M [00:00<00:00, 178MB/s]
 3%|██████████| 1/30 [01:05<31:40, 65.53s/it]
Best Model Saved with Test Loss:  0.43397557735443115
 7%|██████████| 2/30 [01:08<13:26, 28.82s/it]
Best Model Saved with Test Loss:  0.4210205078125
10%|██████████| 3/30 [01:12<07:49, 17.38s/it]
Best Model Saved with Test Loss:  0.34245532751083374
13%|██████████| 4/30 [01:16<05:15, 12.13s/it]
Best Model Saved with Test Loss:  0.20578938722610474
17%|██████████| 5/30 [01:19<03:38, 8.72s/it]
Best Model Saved with Test Loss:  0.1909628709157308
23%|██████████| 7/30 [01:24<02:02, 5.31s/it]
Best Model Saved with Test Loss:  0.1621051033337911
30%|██████████| 9/30 [01:29<01:18, 3.75s/it]
Best Model Saved with Test Loss:  0.14238115151723227
40%|██████████| 12/30 [01:36<00:53, 2.95s/it]
Best Model Saved with Test Loss:  0.1290393869082133
47%|██████████| 14/30 [01:41<00:42, 2.69s/it]
Best Model Saved with Test Loss:  0.12479044993718465
50%|██████████| 15/30 [01:44<00:38, 2.59s/it]
Best Model Saved with Test Loss:  0.11434956391652425

```

```

53%|██████| 16/30 [01:47<00:37, 2.71s/it]
Best Model Saved with Test Loss: 0.11301668485005696
57%|██████| 17/30 [01:49<00:35, 2.72s/it]
Best Model Saved with Test Loss: 0.10837630430857341
60%|██████| 18/30 [01:52<00:32, 2.71s/it]
Best Model Saved with Test Loss: 0.10548104842503865
63%|██████| 19/30 [01:55<00:29, 2.73s/it]
Best Model Saved with Test Loss: 0.09955310821533203
73%|██████| 22/30 [02:02<00:21, 2.65s/it]
Best Model Saved with Test Loss: 0.0928419828414917
80%|██████| 24/30 [02:07<00:15, 2.54s/it]
Best Model Saved with Test Loss: 0.08815499146779378
87%|██████| 26/30 [02:12<00:09, 2.49s/it]
Best Model Saved with Test Loss: 0.08766223986943562
90%|██████| 27/30 [02:15<00:08, 2.69s/it]
Best Model Saved with Test Loss: 0.08430798848470052
100%|██████| 30/30 [02:24<00:00, 4.80s/it]
Best Model Saved with Test Loss: 0.08100565274556477
65 16

100%|██████| 30/30 [01:09<00:00, 2.32s/it]
65 16
100%|██████| 30/30 [01:09<00:00, 2.31s/it]
65 16
100%|██████| 30/30 [01:09<00:00, 2.31s/it]
65 16
100%|██████| 30/30 [01:09<00:00, 2.31s/it]
Cross Validation Results: [0.08100565274556477, 0.09677454829216003, 0.14236590266227722, 0.117516070604

Average_loss: 0.10587898890177408

```

1 # here best test loader means the split for which the test loss was recorded at its minimum.  
 2 len(best\_test\_loader.dataset)

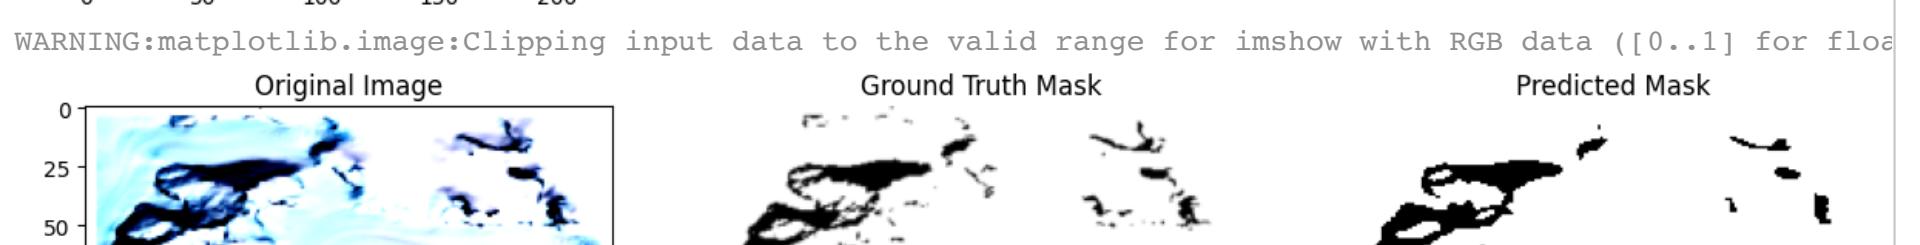
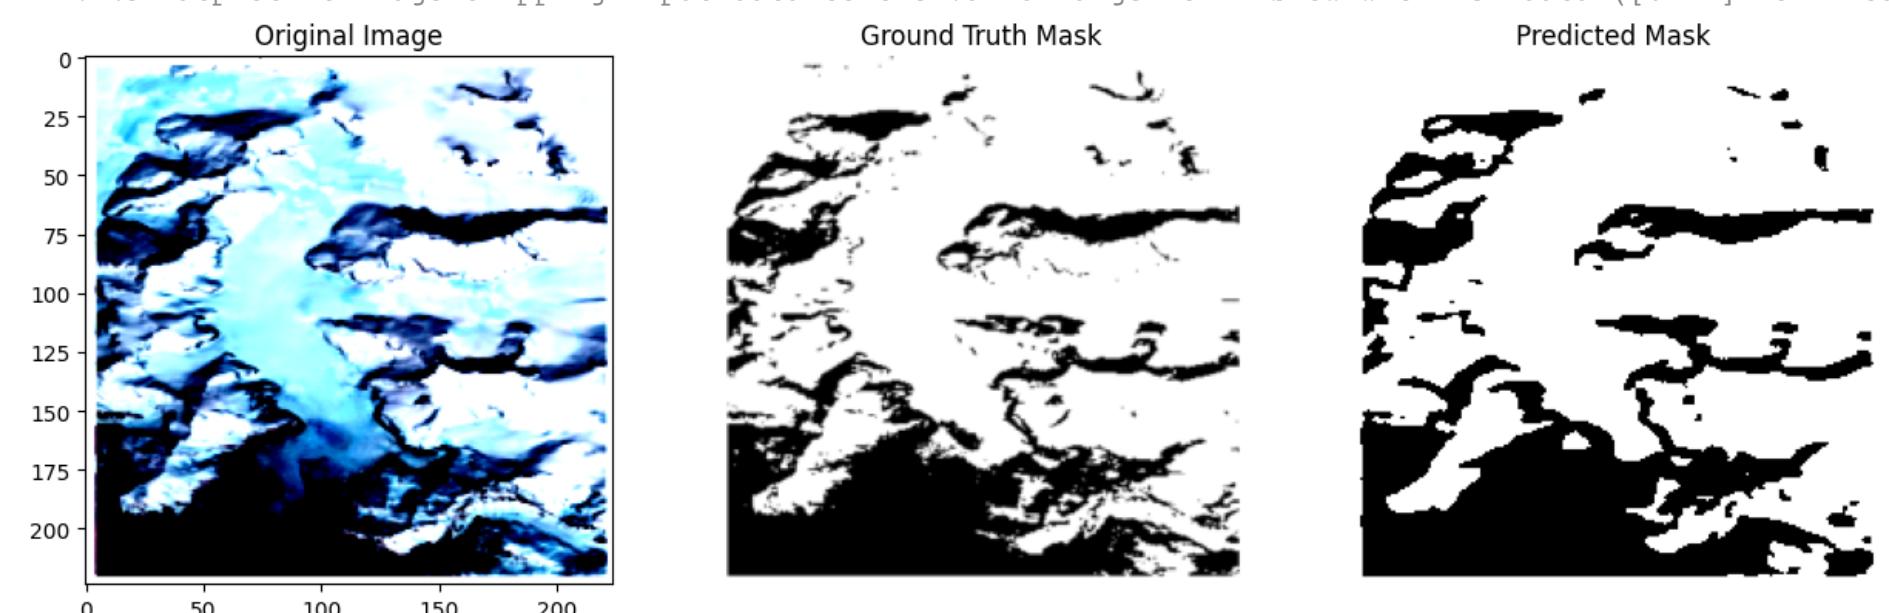
17

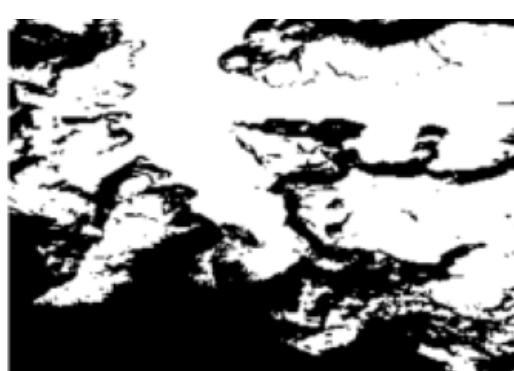
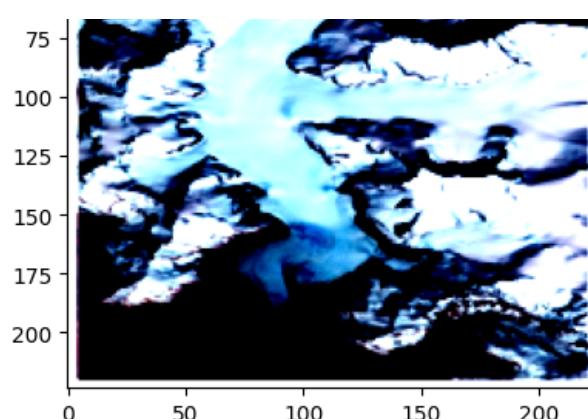
```

1 # Visualizing the model output on the best-picked test data, as seen the predicted mask won't be as
2 # and more training, for this case, i felt that 30 epochs were sufficient to produce acceptable output
3 # However incase of larger region of interest and more classes or cases with instance segmentation,
4 model = smp.Unet(
5     encoder_name="resnet34",
6     encoder_weights="imagenet",
7     in_channels=3,
8     classes=1
9 )
10 model = model.to(device)
11 model.load_state_dict(torch.load('Best_model_resnet34.pth'))
12 test_image_plot(best_test_loader, model)

```

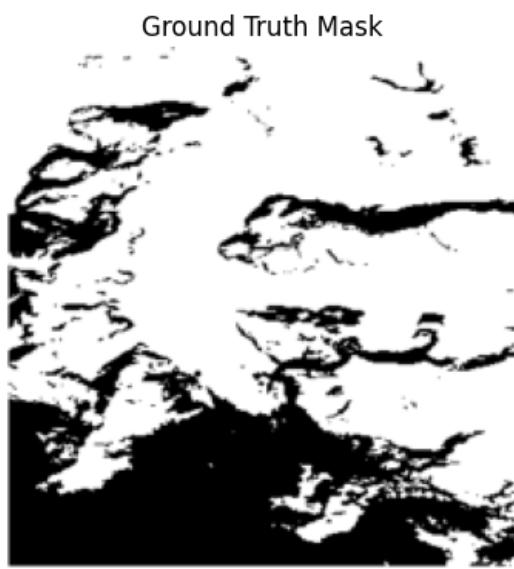
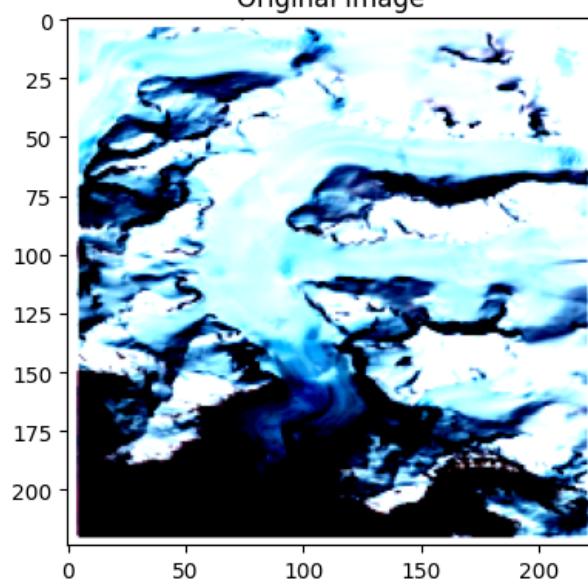
<ipython-input-20-1532a440afdb>:11: FutureWarning: You are using `torch.load` with `weights\_only=False` (model.load\_state\_dict(torch.load('Best\_model\_resnet34.pth')))  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32...





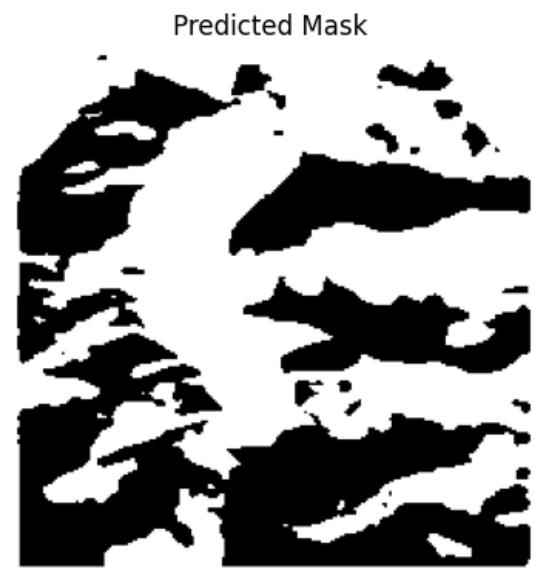
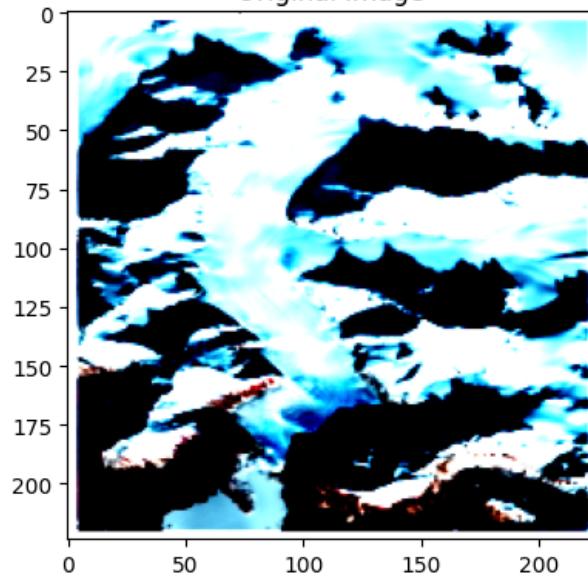
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float type)

Original Image



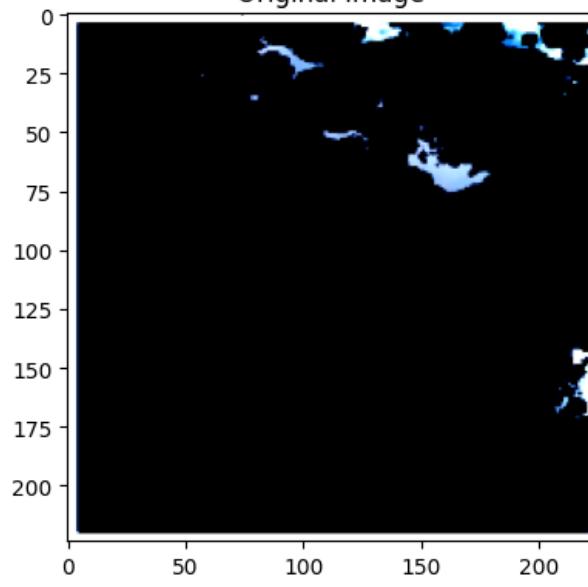
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float type)

Original Image



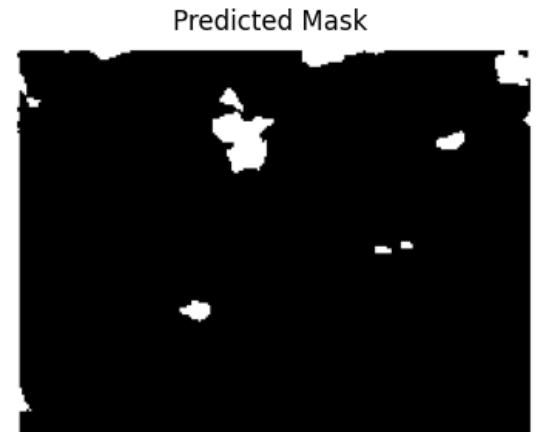
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float type)

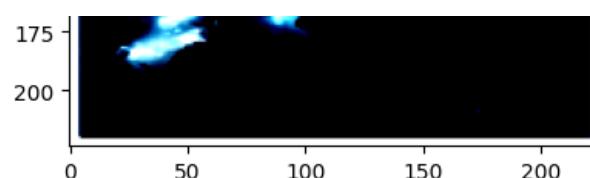
Original Image



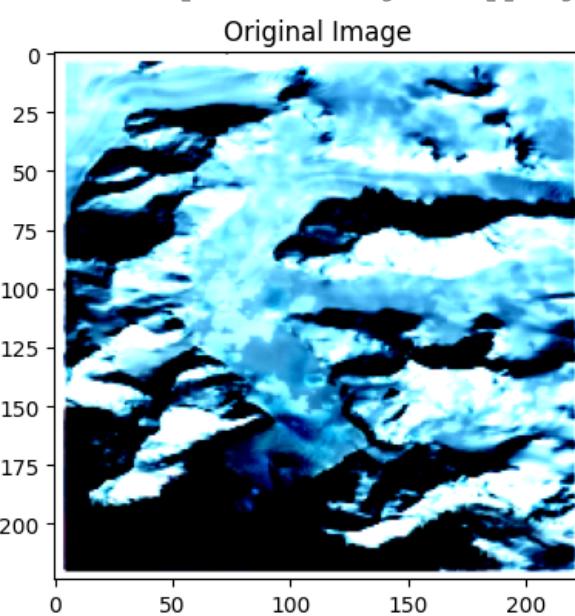
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float type)

Original Image





WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type conversion required

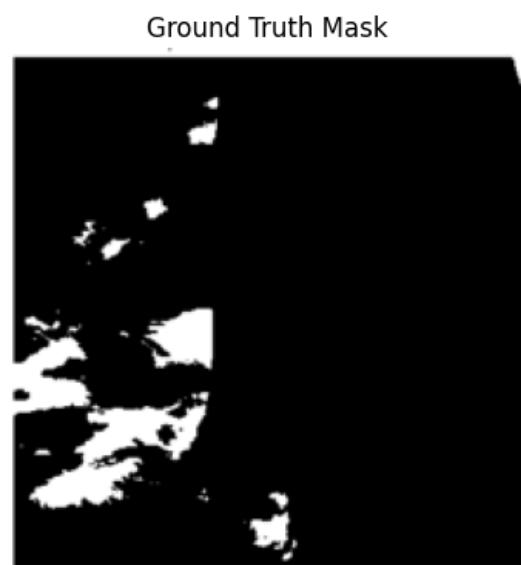
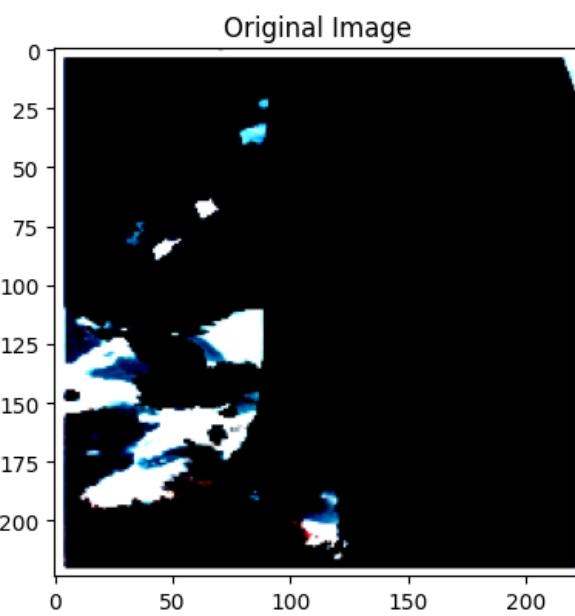


Ground Truth Mask

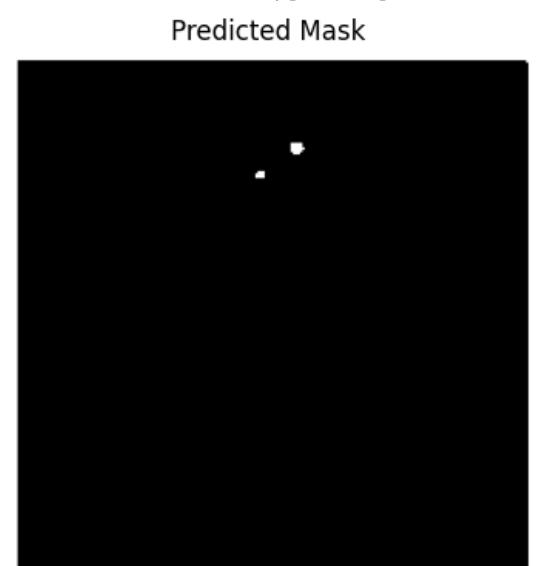
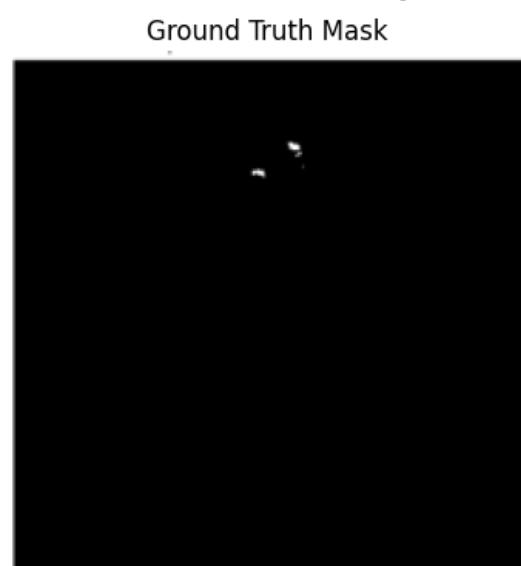
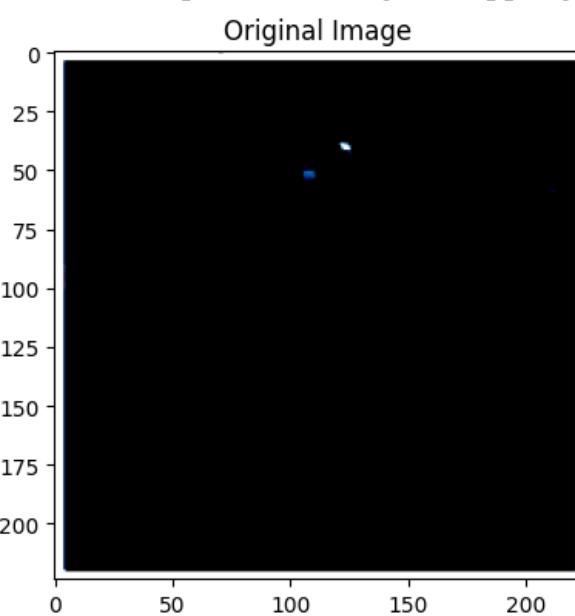


Predicted Mask

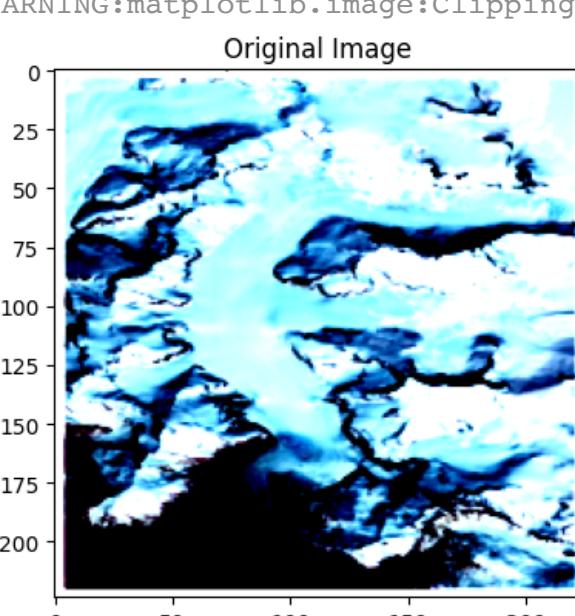
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type conversion required



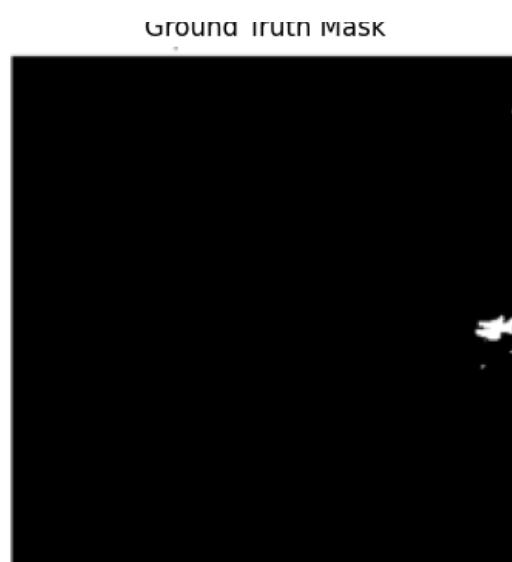
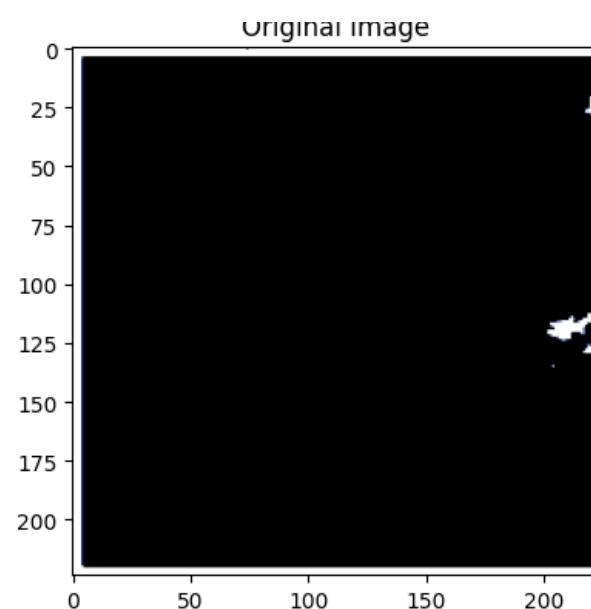
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type conversion required



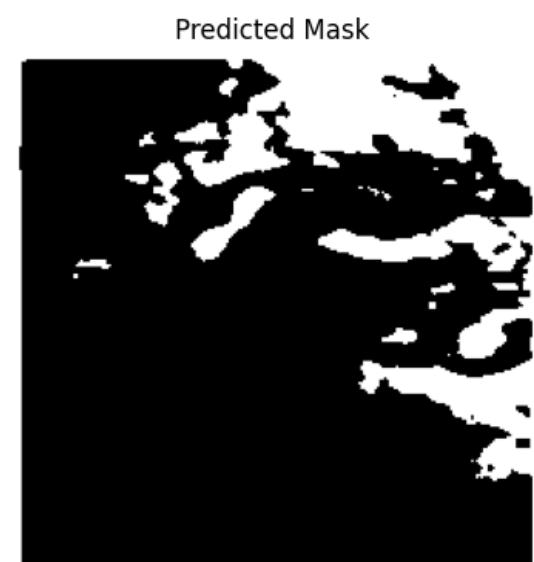
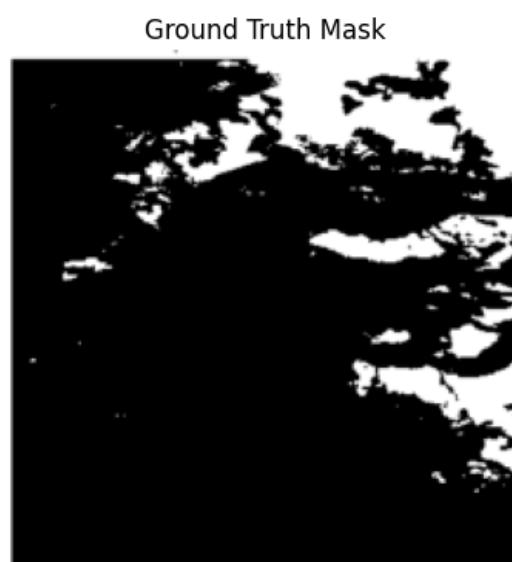
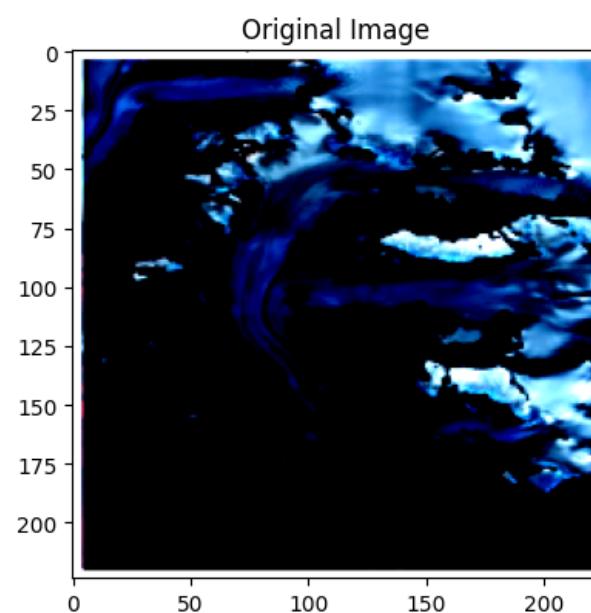
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type conversion required



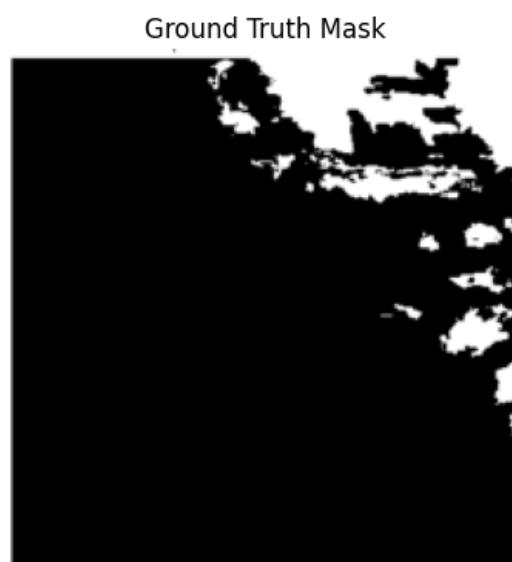
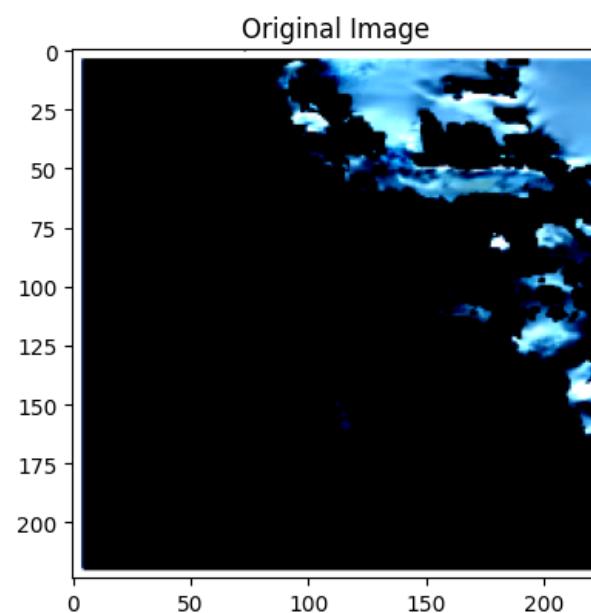
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type conversion required



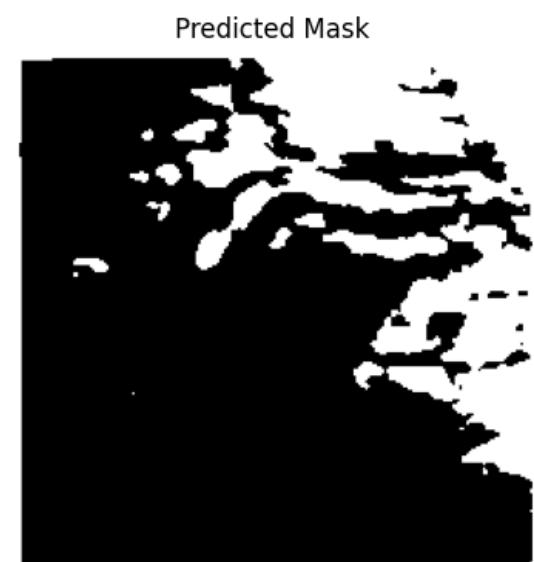
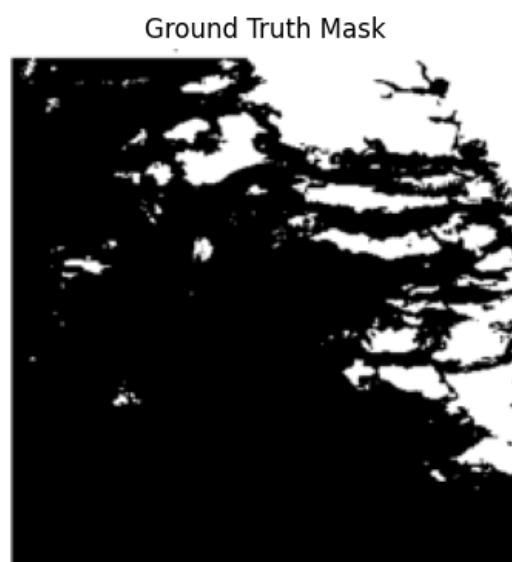
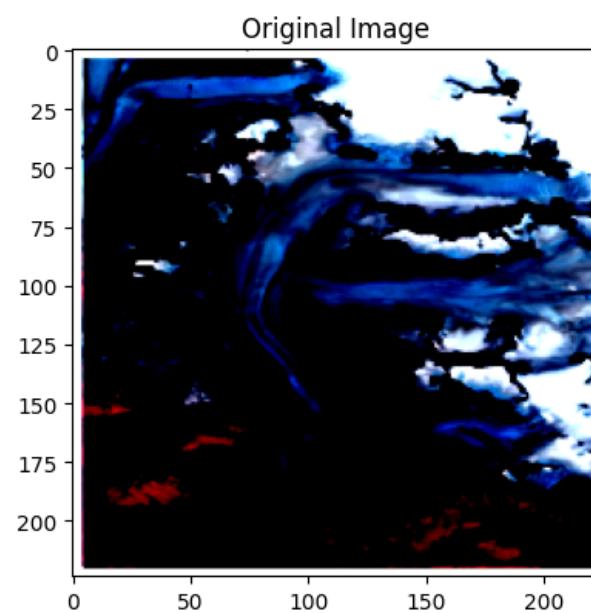
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



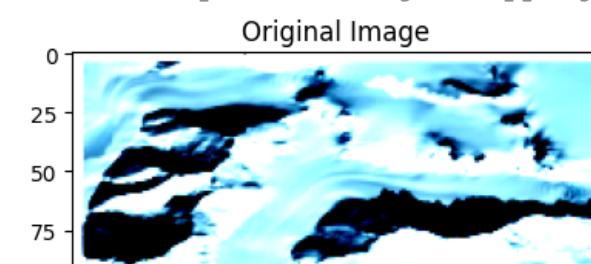
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types

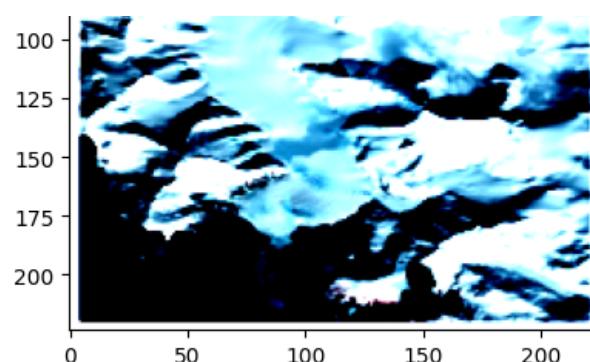


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types

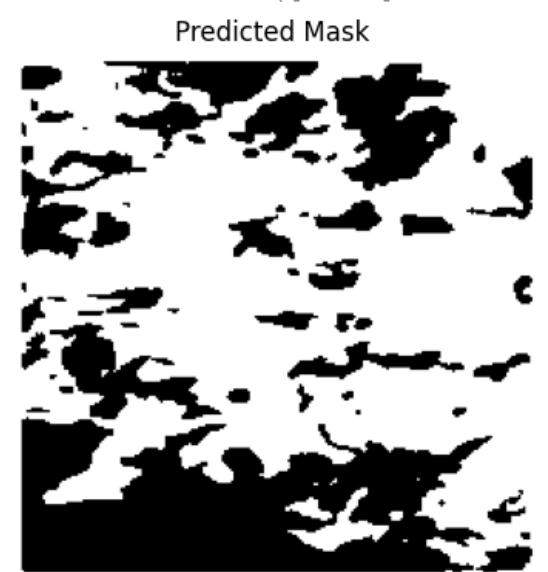
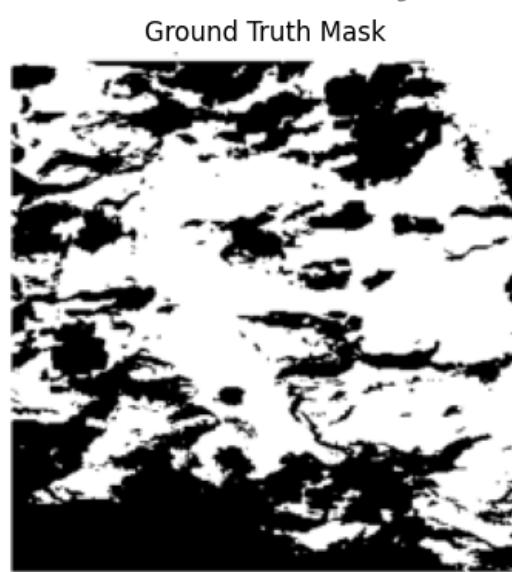
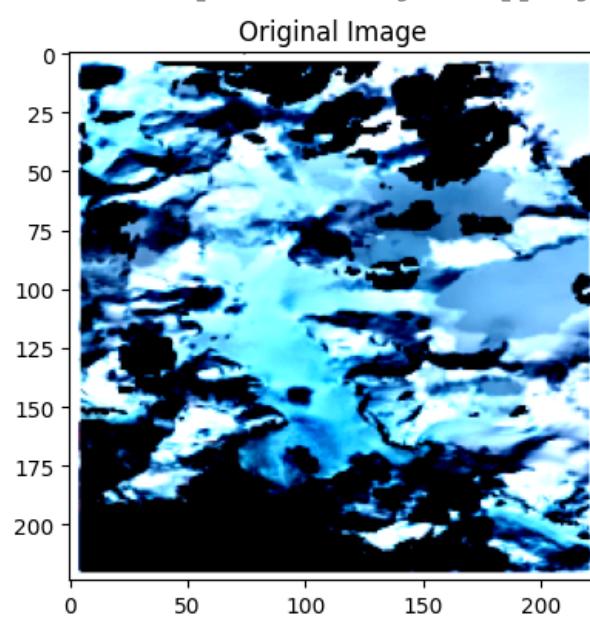


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types

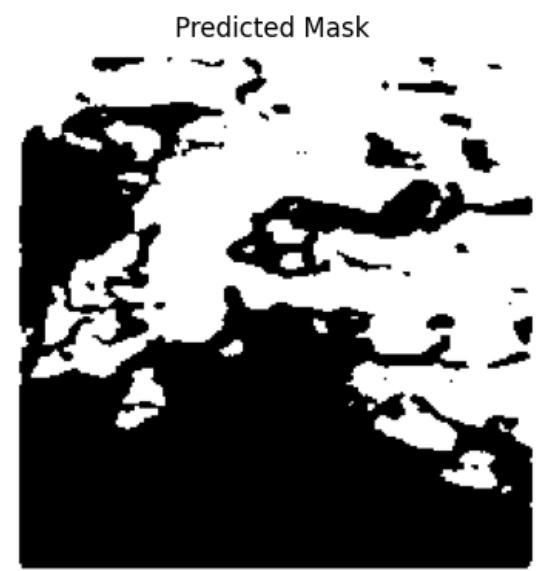
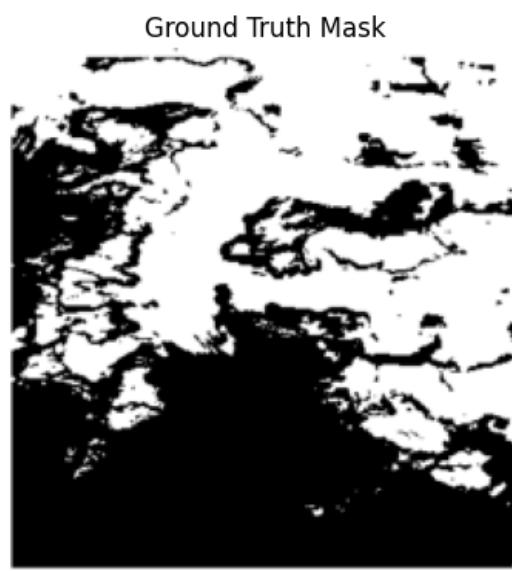
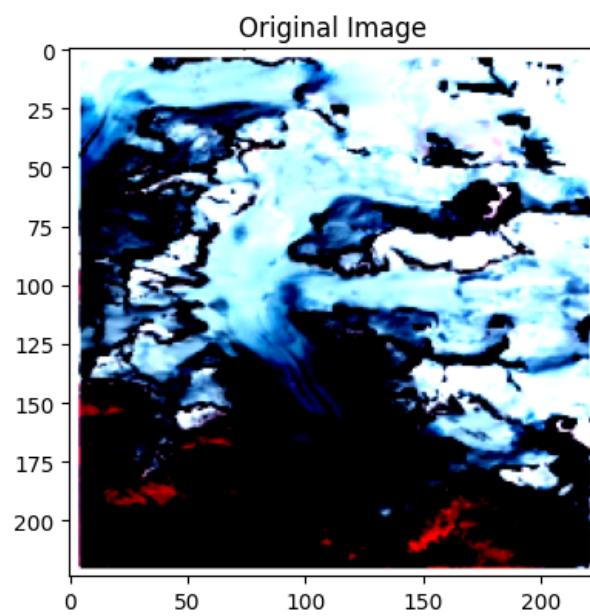




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)



## ✓ UNet with Encoder ResNet50 - Training and Evaluation

```

1 # keeping all the parameters same for fair comparison, except the encoder, this time, I am using Resnet50
2 k_folds_results = []
3 resnet50_best_test_loss = 10000
4 loss = smp.losses.DiceLoss(mode='binary')
5 for fold, (train_index, test_index) in enumerate(kf.split(indices)):
6     train_image_paths = [image_paths[i] for i in train_index]
7     train_mask_paths = [mask_paths[i] for i in train_index]
8     test_image_paths = [image_paths[i] for i in test_index]
9     test_mask_paths = [mask_paths[i] for i in test_index]
10    print(len(train_index), len(test_index))
11    train_dataset = SegmentationDataset(train_image_paths, train_mask_paths)
12    test_dataset = SegmentationDataset(test_image_paths, test_mask_paths)
13    train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
14    test_loader = DataLoader(test_dataset, batch_size=8, shuffle=True)
15    model = smp.Unet(
16        encoder_name="resnet50",
17        encoder_weights="imagenet",
18        in_channels=3,
19        classes=1
20    )
21    model = model.to(device)
22    optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)
23    for epoch in tqdm(range(30)):
24        train_loss = train(model, train_loader, optimizer, loss, device)
25        test_loss = test(model, test_loader, loss, device)
26        if test_loss < resnet50_best_test_loss:
27            resnet50_best_test_loss = test_loss
28            torch.save(model.state_dict(), f"Best_model_resnet50.pth")
29            print("\nBest Model Saved with Test Loss: ", resnet50_best_test_loss)
30            resnet50_best_test_loader = test_loader
31    k_folds_results.append(test_loss)
32 print("\nCross Validation Results: ", k_folds_results)
33 print("\nAverage loss: ", sum(k_folds_results)/len(k_folds_results))

```

```

64 17
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.cache/torch/hub/checkpoints/resnet50-19c8e357.pth
100%[██████████] 97.8M/97.8M [00:00<00:00, 216MB/s]
 3%[██]          | 1/30 [00:03<01:34,  3.25s/it]
Best Model Saved with Test Loss:  0.5683464209238688
 7%[██]          | 2/30 [00:06<01:29,  3.20s/it]
Best Model Saved with Test Loss:  0.31959615151087445
13%[███]         | 4/30 [00:12<01:19,  3.05s/it]
Best Model Saved with Test Loss:  0.20027210315068564
30%[███]         | 9/30 [00:26<01:01,  2.91s/it]
Best Model Saved with Test Loss:  0.1463657021522522
33%[███]         | 10/30 [00:30<01:00,  3.03s/it]
Best Model Saved with Test Loss:  0.1407331625620524
37%[███]         | 11/30 [00:33<00:58,  3.06s/it]
Best Model Saved with Test Loss:  0.14008078972498575
43%[███]         | 13/30 [00:39<00:52,  3.10s/it]
Best Model Saved with Test Loss:  0.1374316414197286
50%[███]         | 15/30 [00:45<00:44,  2.97s/it]
Best Model Saved with Test Loss:  0.12121532360712688
53%[███]         | 16/30 [00:48<00:44,  3.19s/it]
Best Model Saved with Test Loss:  0.11623650789260864
60%[███]         | 18/30 [00:55<00:37,  3.13s/it]
Best Model Saved with Test Loss:  0.11049691836039226
63%[███]         | 19/30 [00:58<00:34,  3.15s/it]
Best Model Saved with Test Loss:  0.10581533114115398
70%[███]         | 21/30 [01:04<00:28,  3.19s/it]
Best Model Saved with Test Loss:  0.10520986715952556
83%[███]         | 25/30 [01:16<00:14,  2.98s/it]
Best Model Saved with Test Loss:  0.09534839789072673
87%[███]         | 26/30 [01:19<00:12,  3.15s/it]
Best Model Saved with Test Loss:  0.09269265333811443
93%[███]         | 28/30 [01:25<00:06,  3.07s/it]
Best Model Saved with Test Loss:  0.09195496638615926
100%[███]        | 30/30 [01:31<00:00,  3.06s/it]
Best Model Saved with Test Loss:  0.08649138609568278
65 16
27%[██]          | 8/30 [00:21<00:59,  2.70s/it]

```

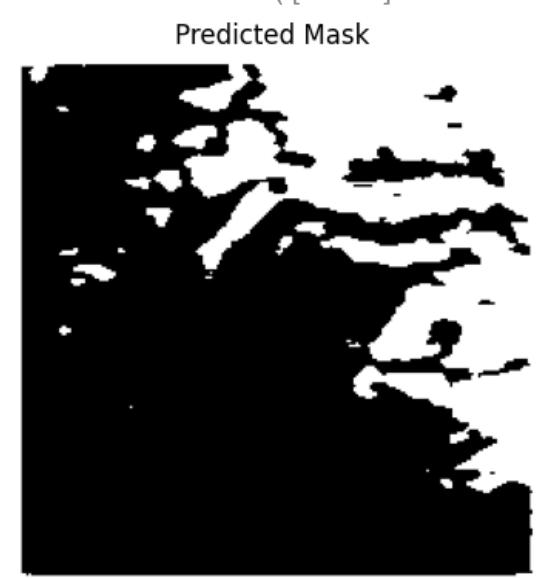
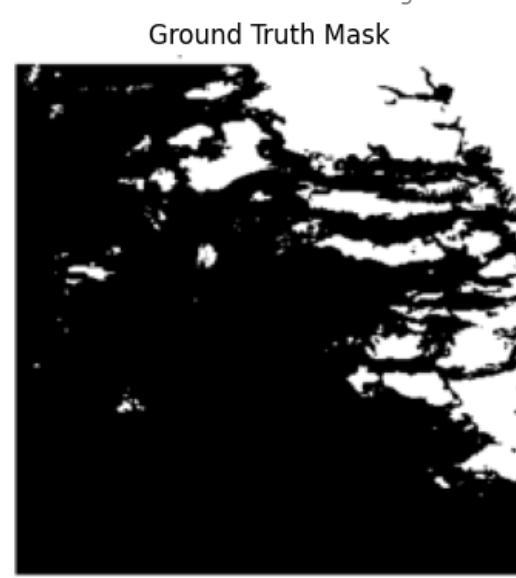
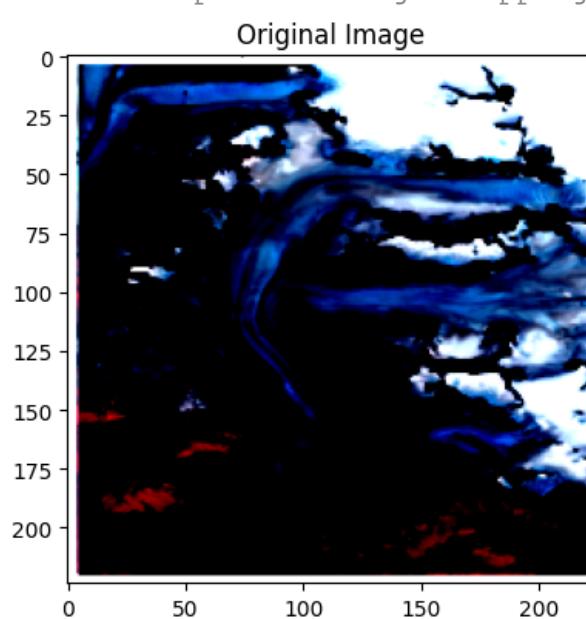
```

1 # visualizing the segmented output, similarly, on the best recorded test split.

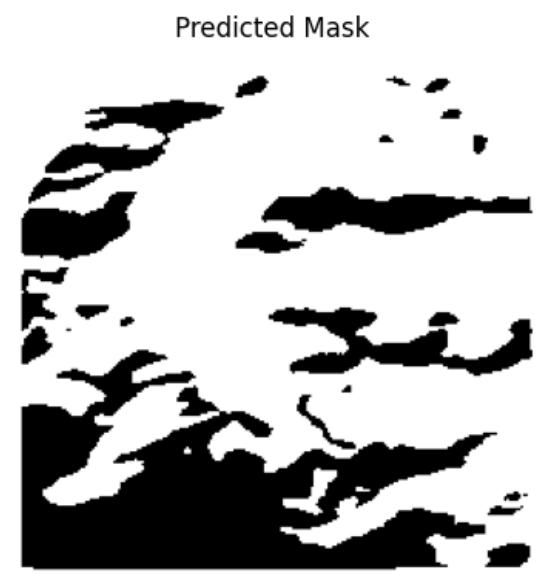
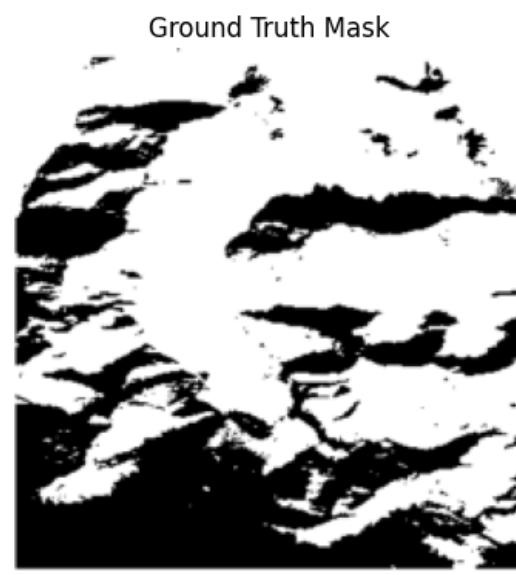
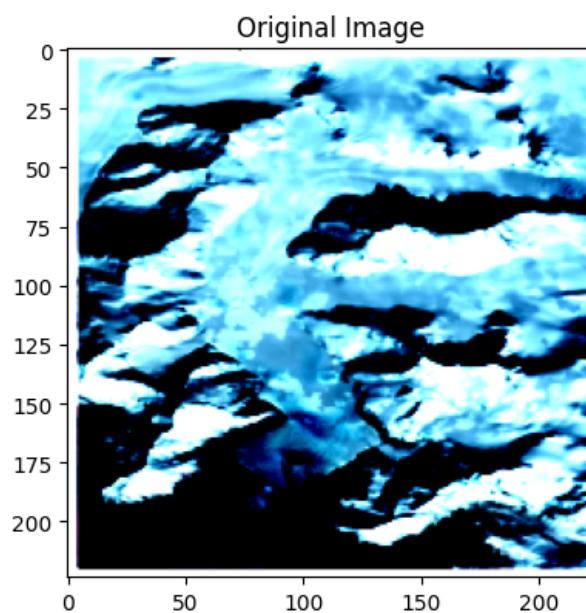
```

```
2 model = smp.Unet(  
3     encoder_name="resnet50",  
4     encoder_weights="imagenet",  
5     in_channels=3,  
6     classes=1  
7 )  
8 model = model.to(device)  
9 model.load_state_dict(torch.load('Best_model_resnet50.pth'))  
10 test_image_plot(resnet50_best_test_loader, model)
```

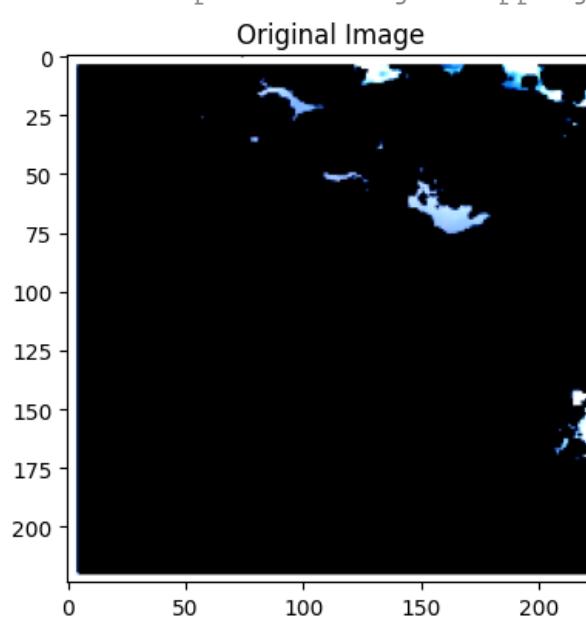
<ipython-input-22-01c725dc0aa7>:9: FutureWarning: You are using `torch.load` with `weights\_only=False` (t  
model.load\_state\_dict(torch.load('Best\_model\_resnet50.pth'))  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floa



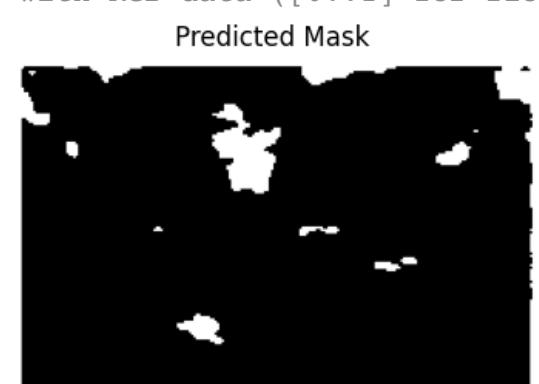
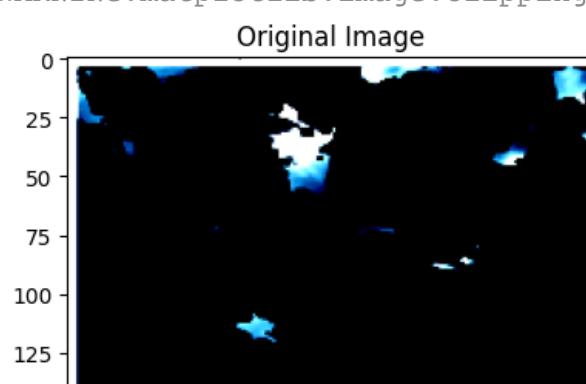
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floa

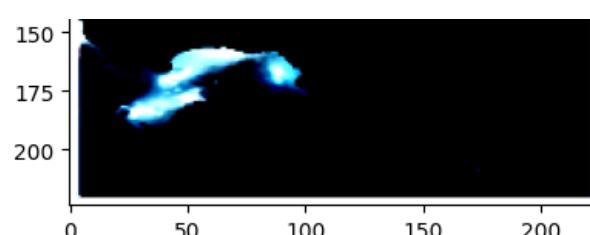


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floa

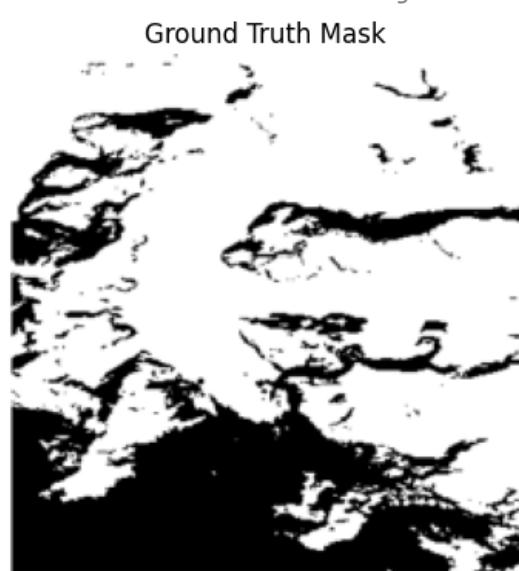
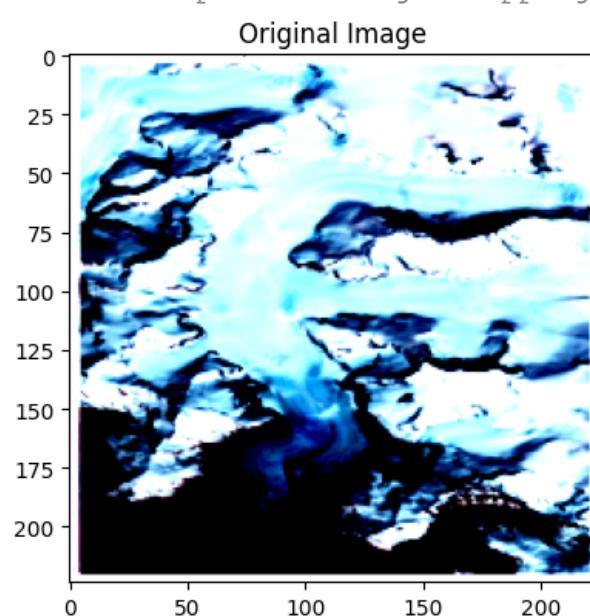


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floa

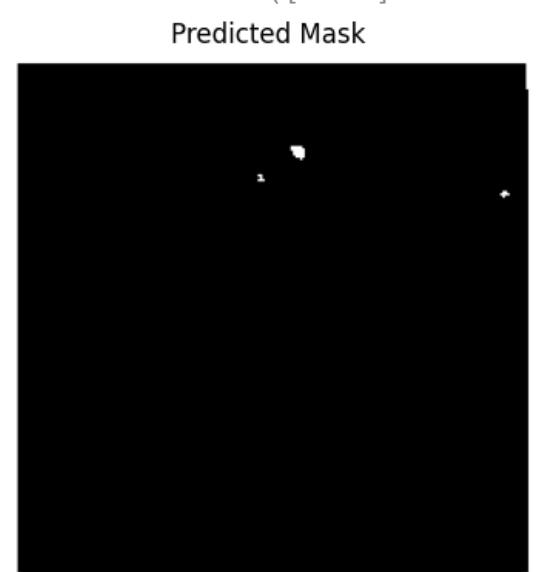
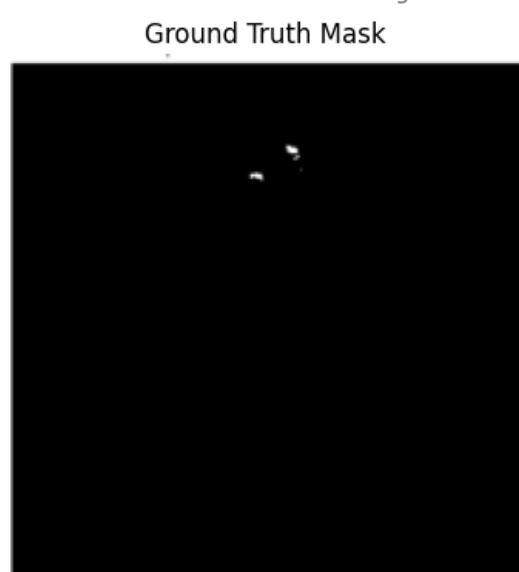
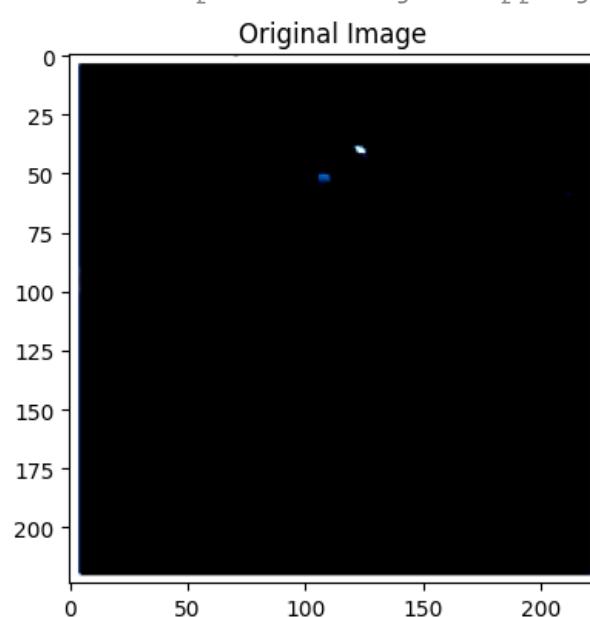




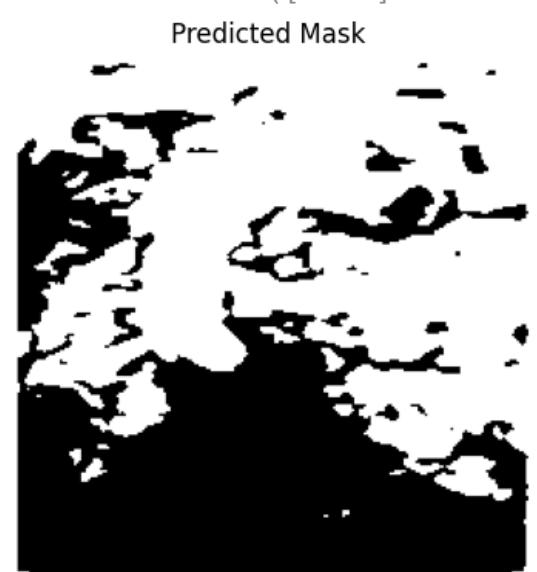
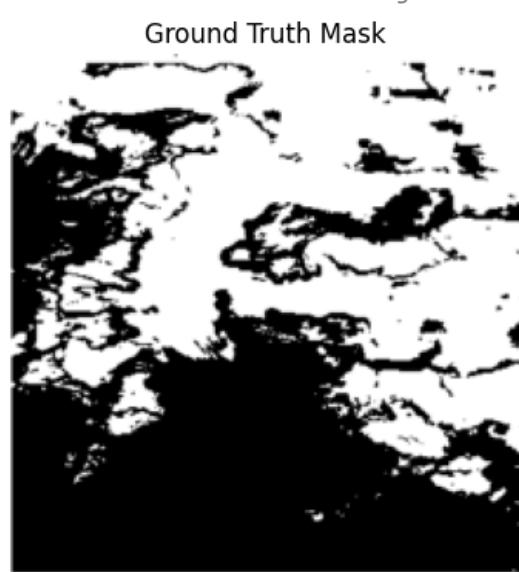
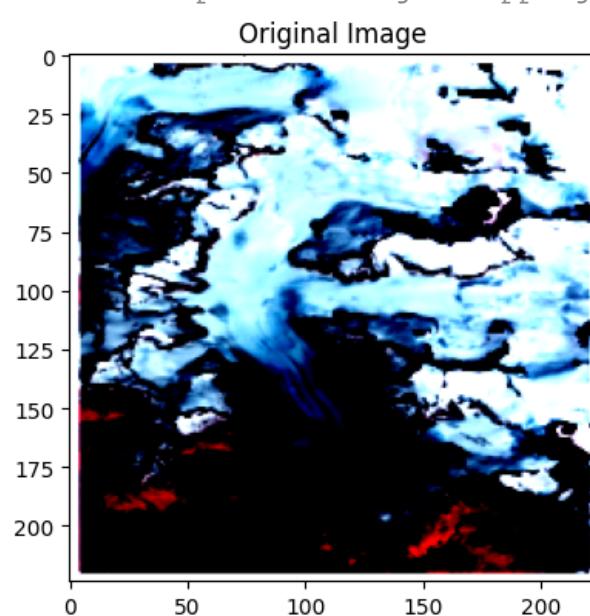
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



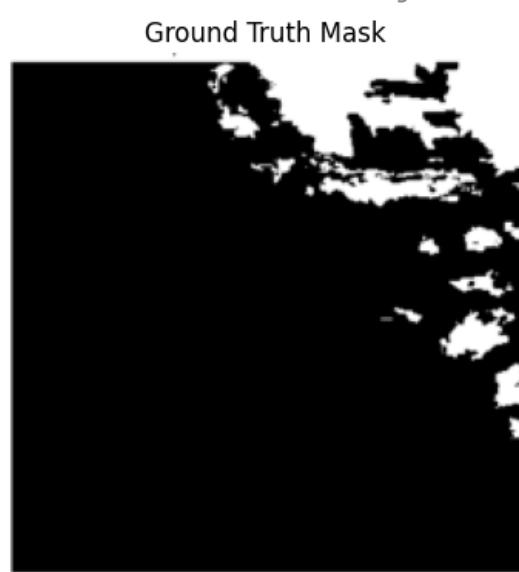
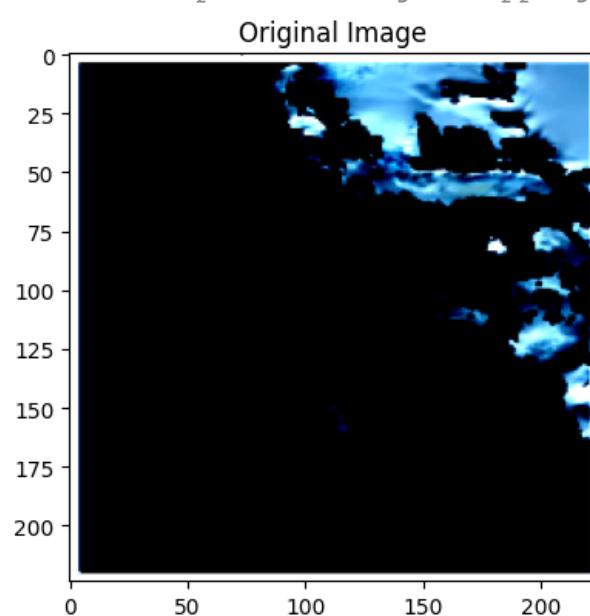
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



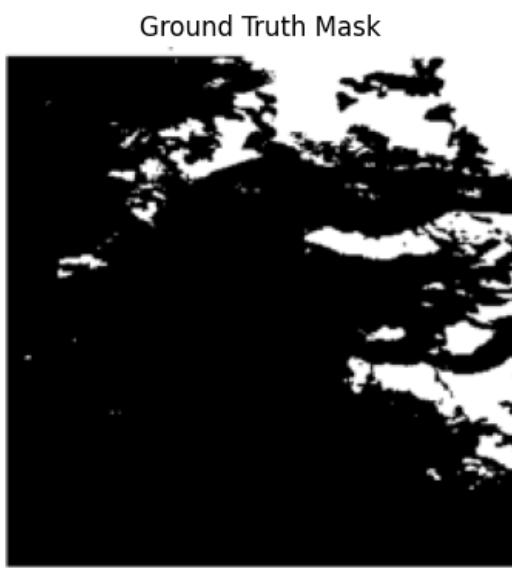
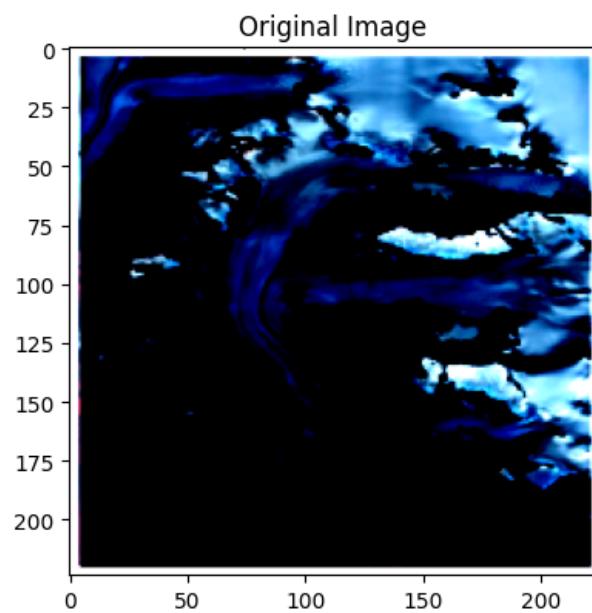
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



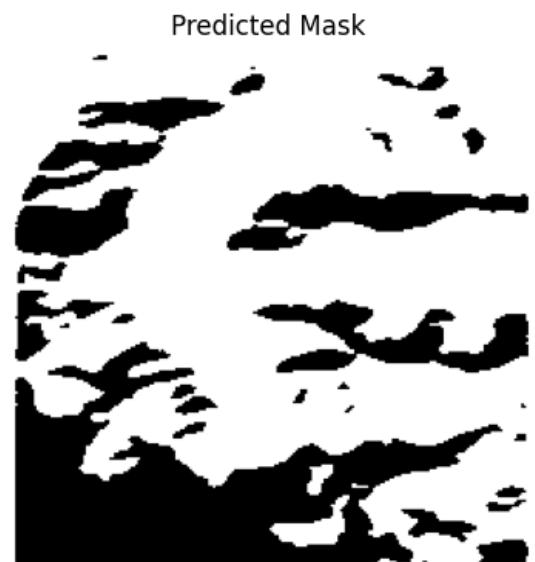
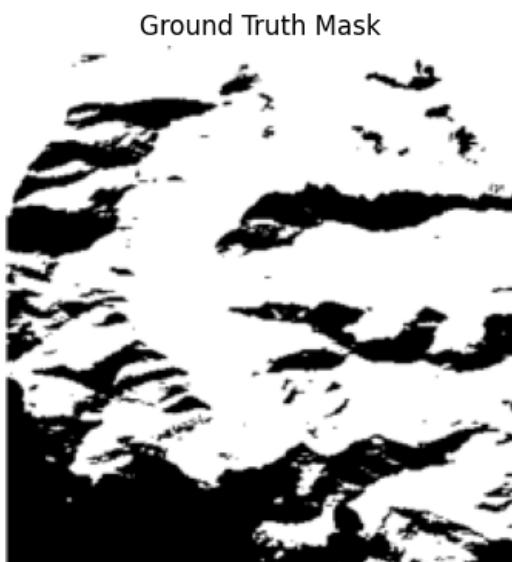
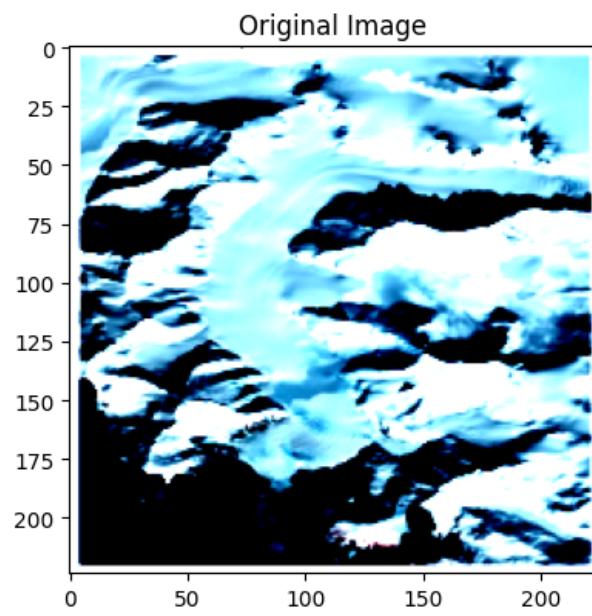
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



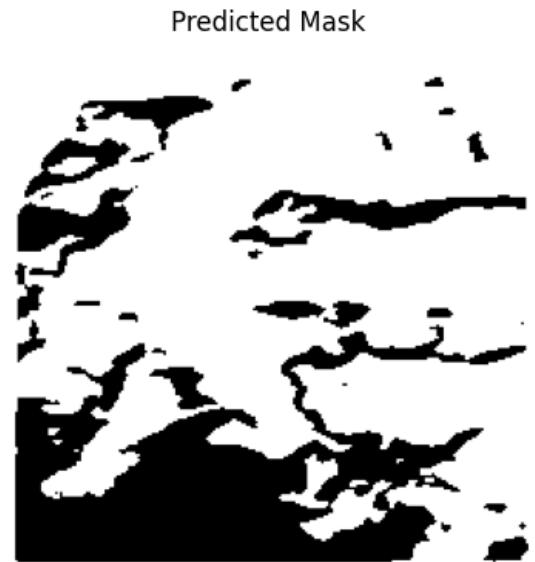
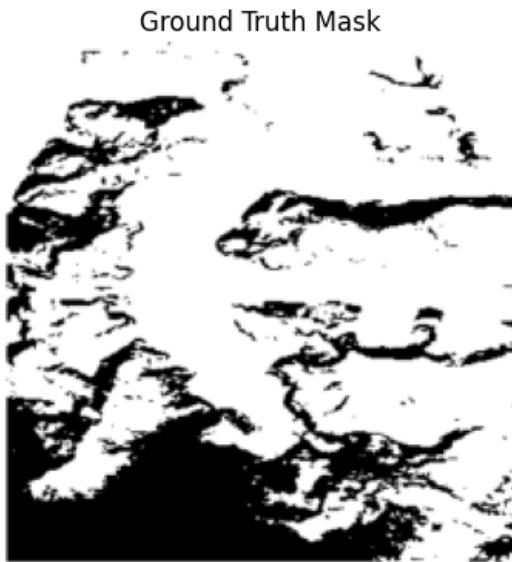
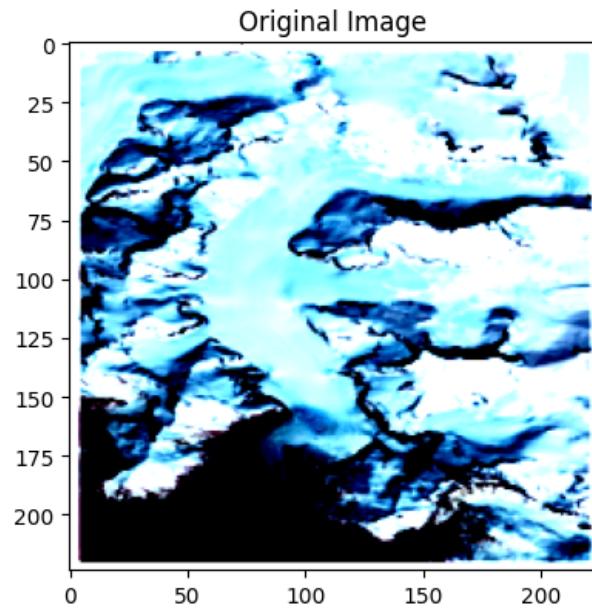
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) to [0..255] for uint8



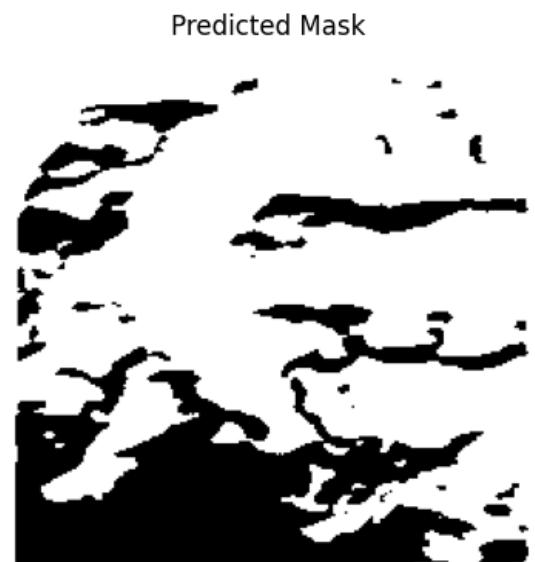
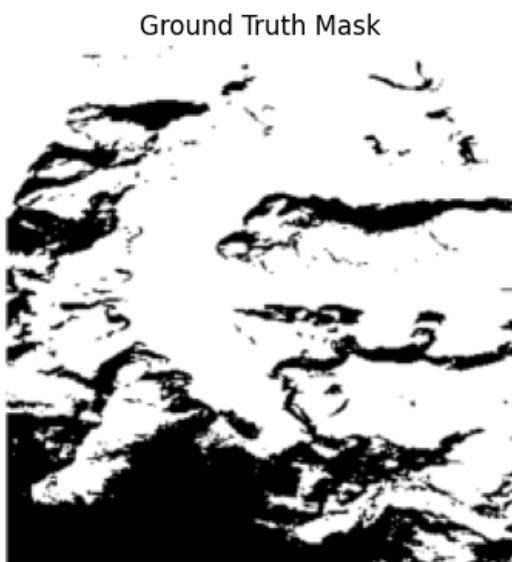
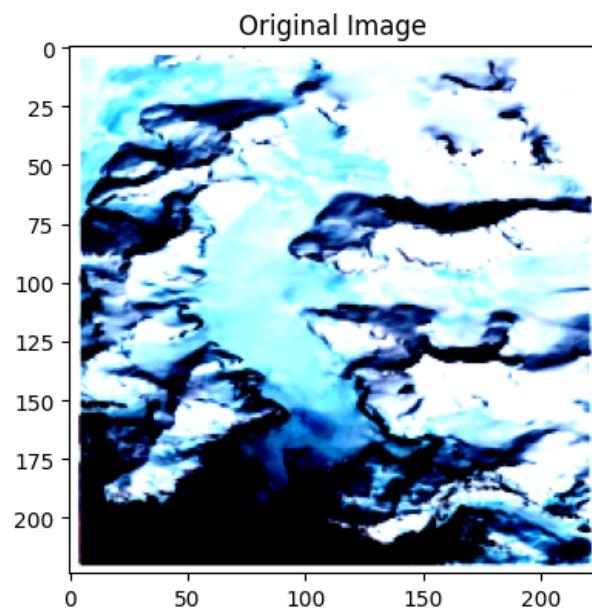
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) to [0..255] for uint8



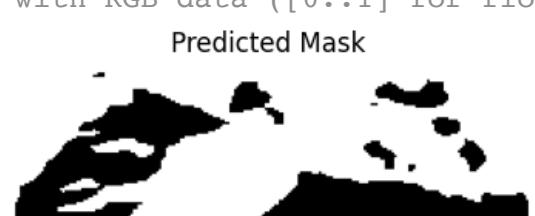
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) to [0..255] for uint8

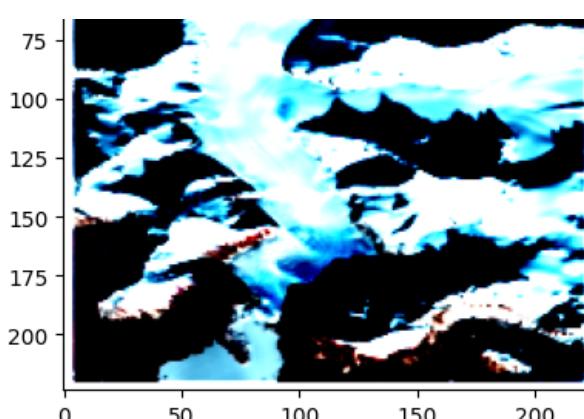


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) to [0..255] for uint8



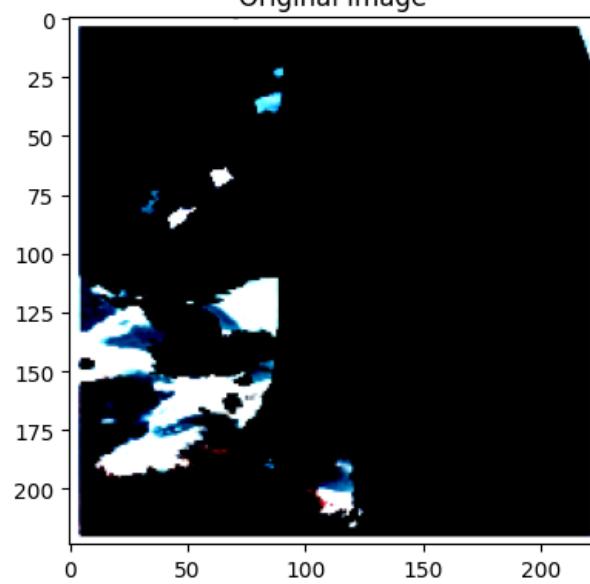
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) to [0..255] for uint8





WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

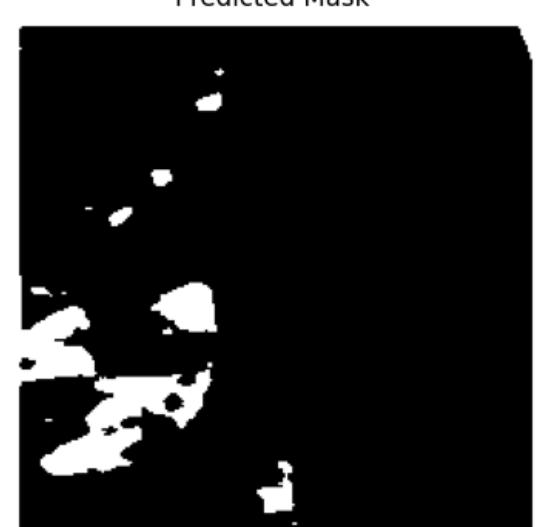
Original Image



Ground Truth Mask

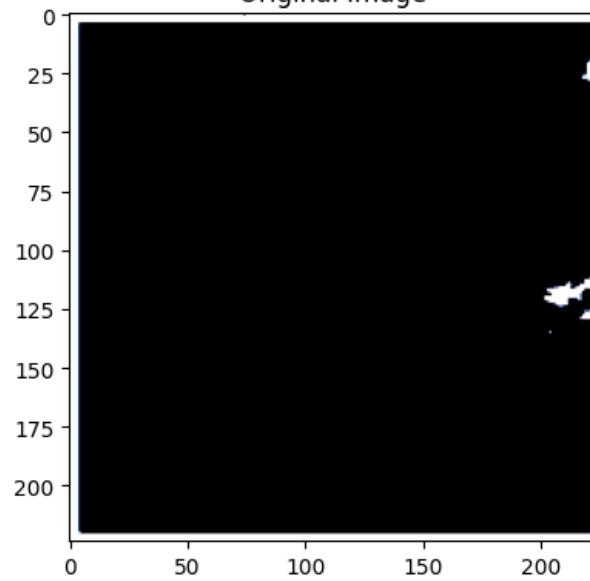


Predicted Mask



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

Original Image



Ground Truth Mask

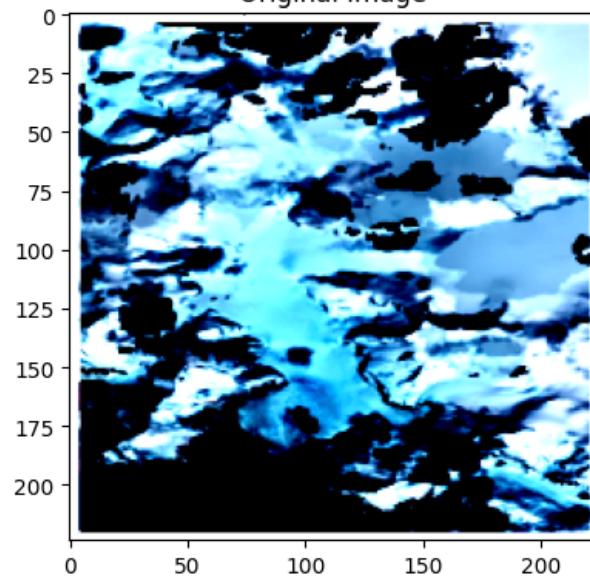


Predicted Mask

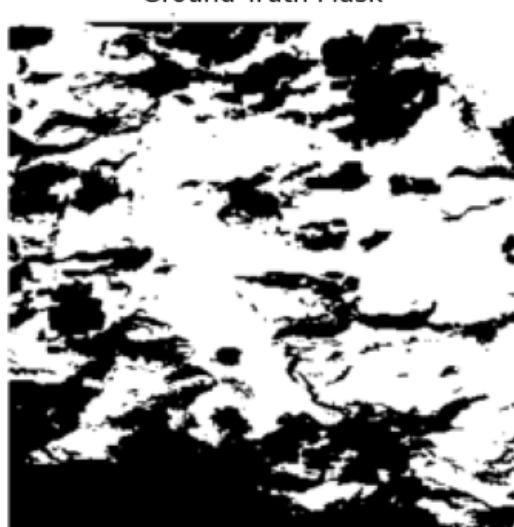


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

Original Image



Ground Truth Mask

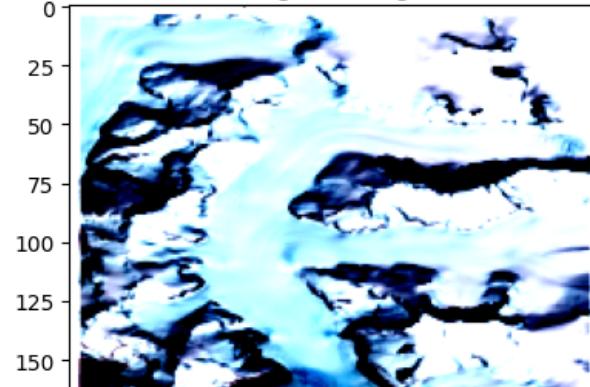


Predicted Mask

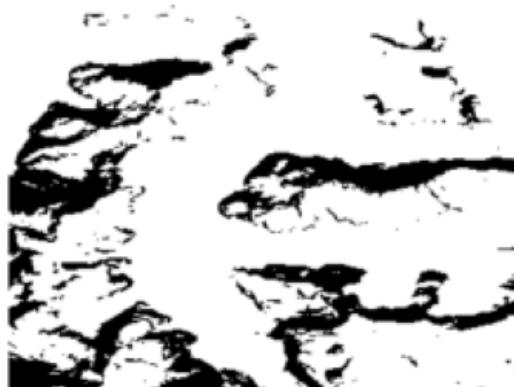


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float data)

Original Image

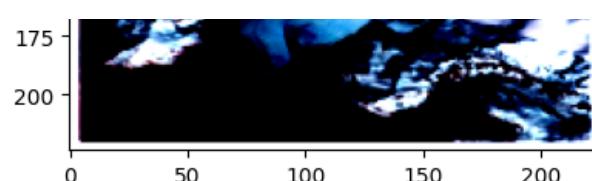


Ground Truth Mask



Predicted Mask





## UNet with Encoder EfficientNet-b4: Training and Evaluation

```

1 # e# keeping all the parameters same for fair comparison, except the encoder, this time, I am using
2 k_folds_results = []
3 efficientb4_best_test_loss = 10000
4 loss = smp.losses.DiceLoss(mode='binary')
5 for fold, (train_index, test_index) in enumerate(kf.split(indices)):
6     train_image_paths = [image_paths[i] for i in train_index]
7     train_mask_paths = [mask_paths[i] for i in train_index]
8     test_image_paths = [image_paths[i] for i in test_index]
9     test_mask_paths = [mask_paths[i] for i in test_index]
10    print(len(train_index), len(test_index))
11    train_dataset = SegmentationDataset(train_image_paths, train_mask_paths)
12    test_dataset = SegmentationDataset(test_image_paths, test_mask_paths)
13    train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
14    test_loader = DataLoader(test_dataset, batch_size=8, shuffle=True)
15    model = smp.Unet(
16        encoder_name="efficientnet-b4",
17        encoder_weights="imagenet",
18        in_channels=3,
19        classes=1
20    )
21    model = model.to(device)
22    optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)
23    for epoch in tqdm(range(30)):
24        train_loss = train(model, train_loader, optimizer, loss, device)
25        test_loss = test(model, test_loader, loss, device)
26        if test_loss < efficientb4_best_test_loss:
27            efficientb4_best_test_loss = test_loss
28            torch.save(model.state_dict(), f"Best_model_efficientnet-b4.pth")
29            print("\nBest Model Saved with Test Loss: ", efficientb4_best_test_loss)
30            best_test_loader = test_loader
31    k_folds_results.append(test_loss)
32 print("\nCross Validation Results: ", k_folds_results)
33 print("\nAverage loss: ", sum(k_folds_results)/len(k_folds_results))

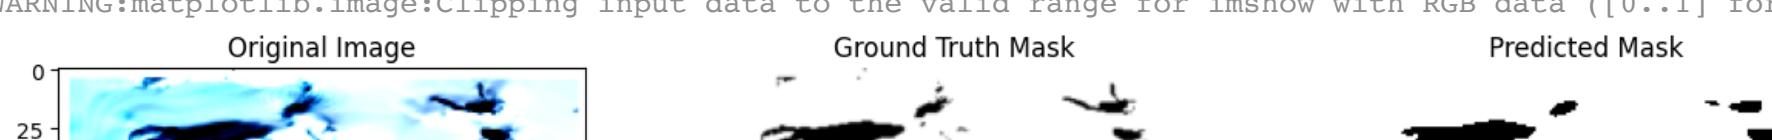
```

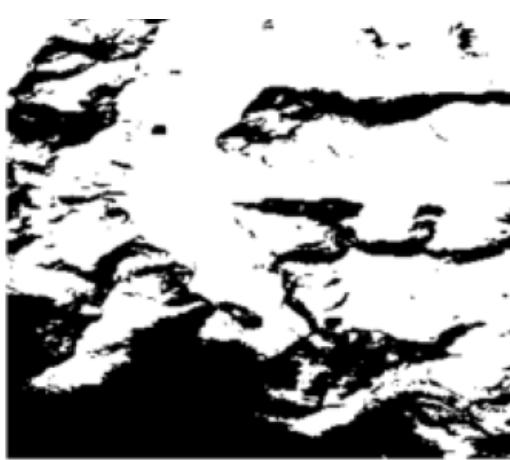
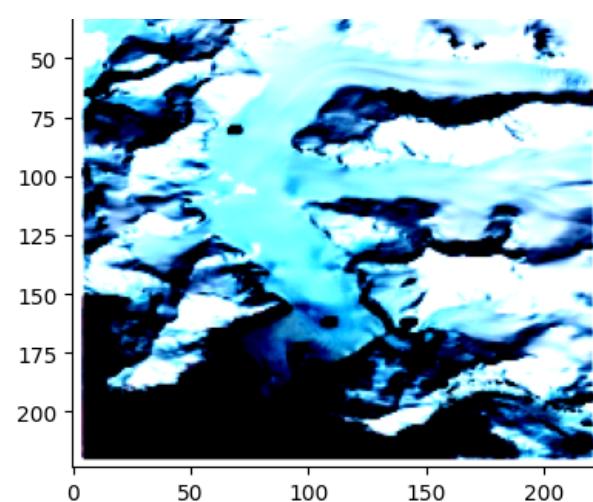
```

1 # again doing the visualization of model performance on segmented mask, using the best test split, I
2 model = smp.Unet(
3     encoder_name="efficientnet-b4",
4     encoder_weights="imagenet",
5     in_channels=3,
6     classes=1
7 )
8 model = model.to(device)
9 model.load_state_dict(torch.load('Best_model_efficientnet-b4.pth'))
10 test_image_plot(best_test_loader, model)

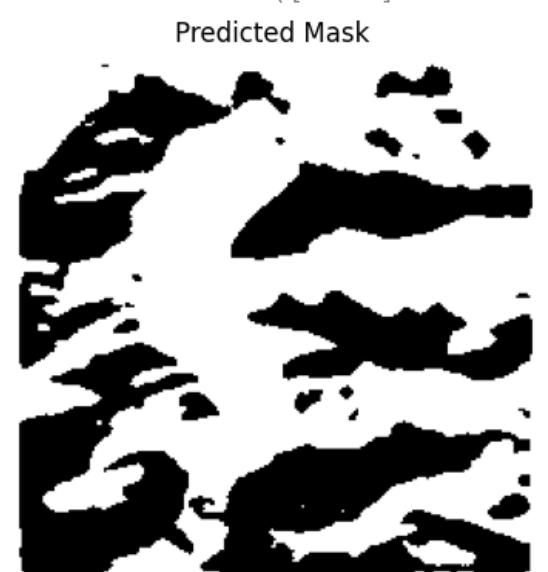
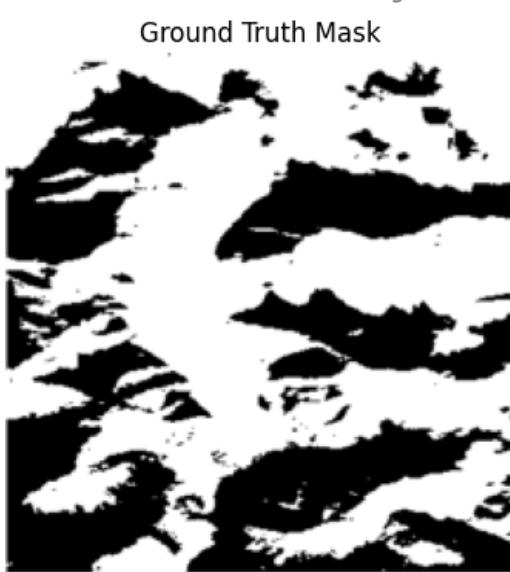
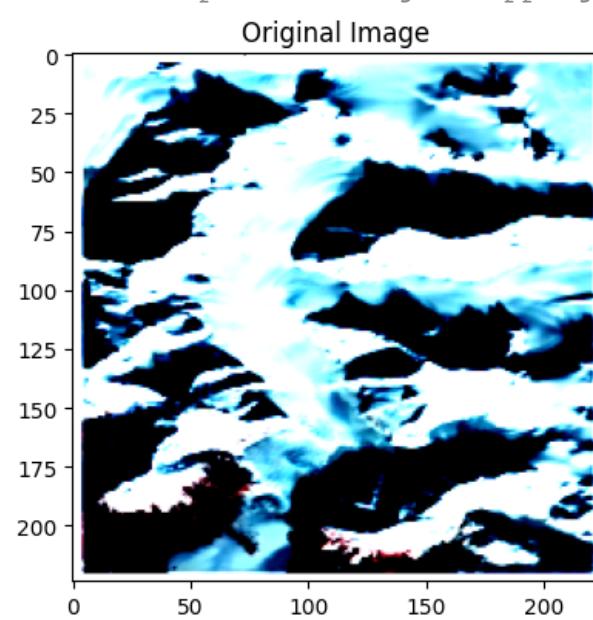
```

<ipython-input-24-d3efdd47aea9>:9: FutureWarning: You are using `torch.load` with `weights\_only=False` (t  
model.load\_state\_dict(torch.load('Best\_model\_efficientnet-b4.pth'))  
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for flo

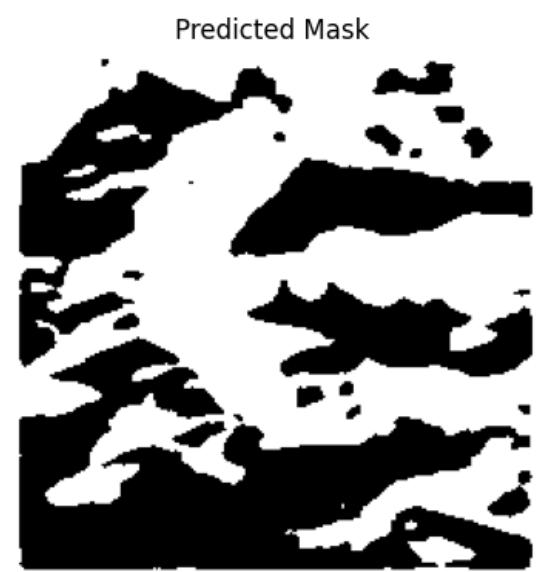
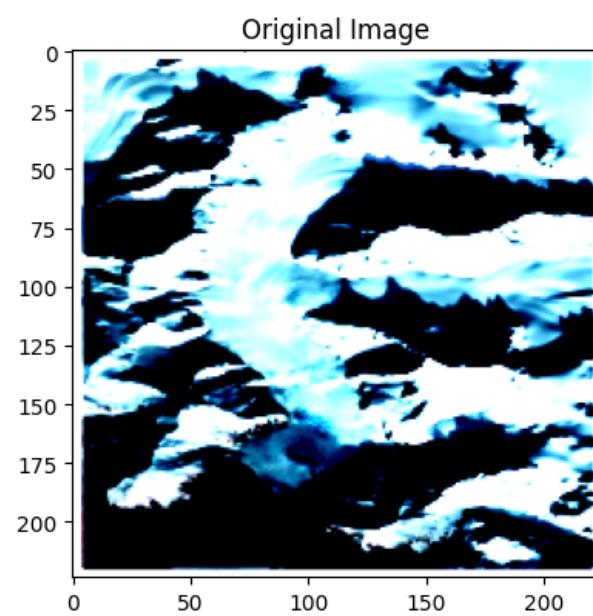




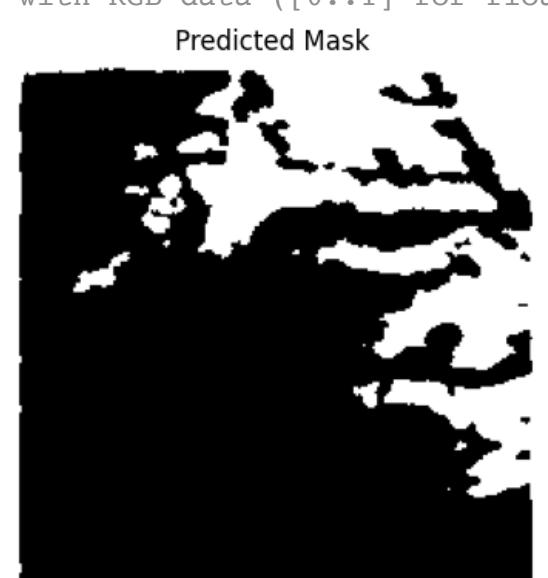
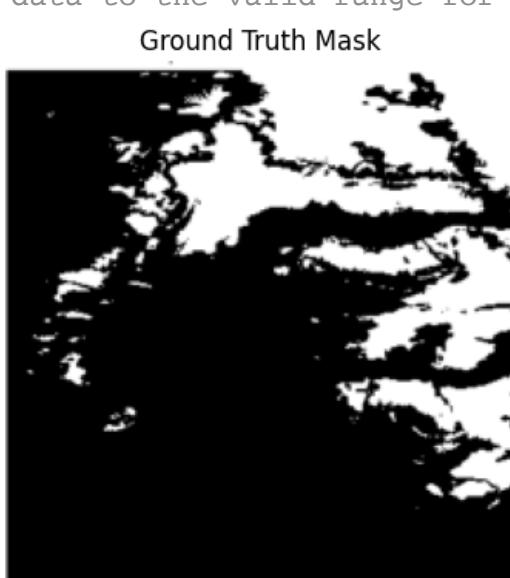
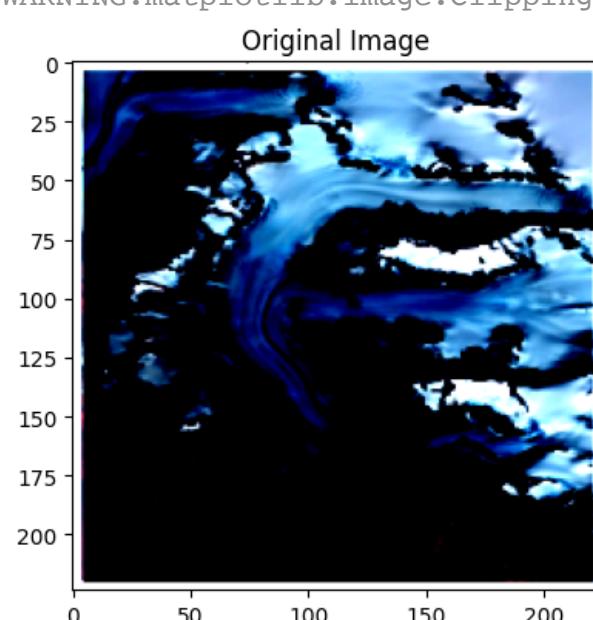
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar)



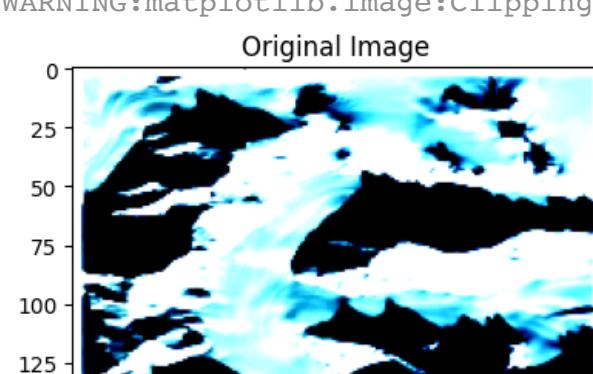
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar)

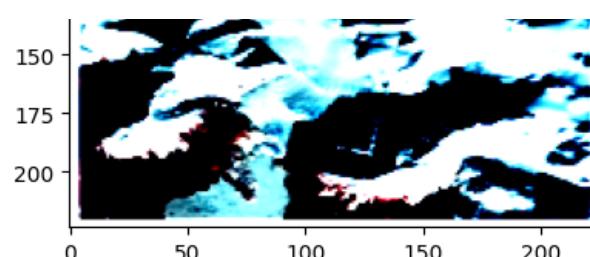


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar)

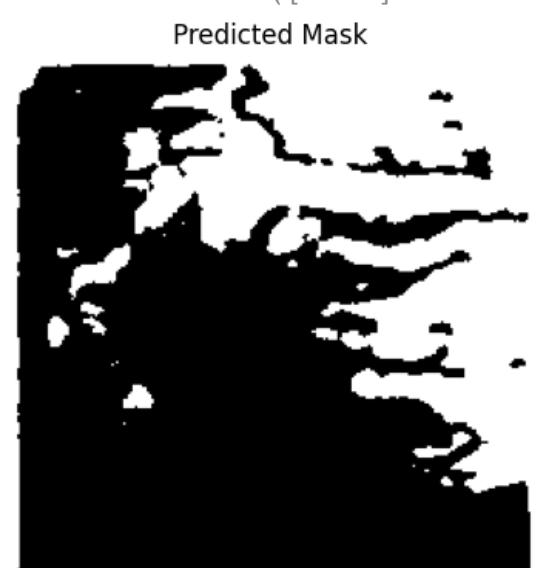
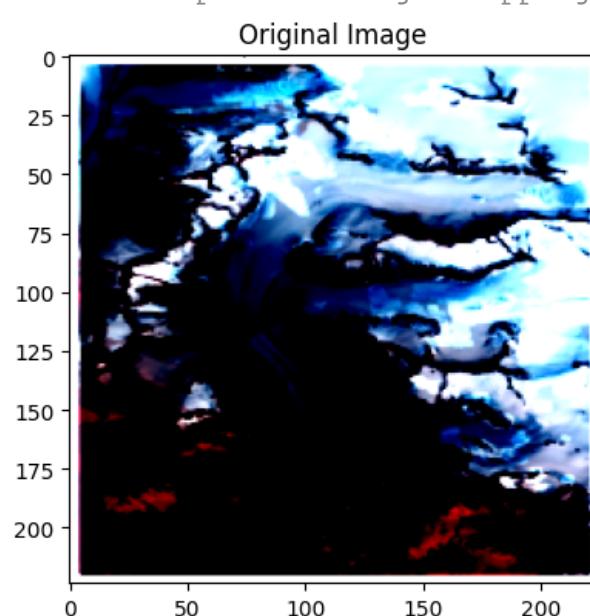


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar)

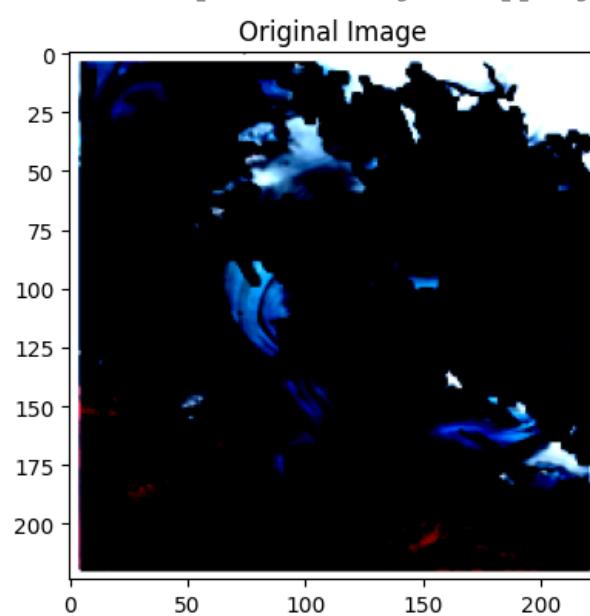




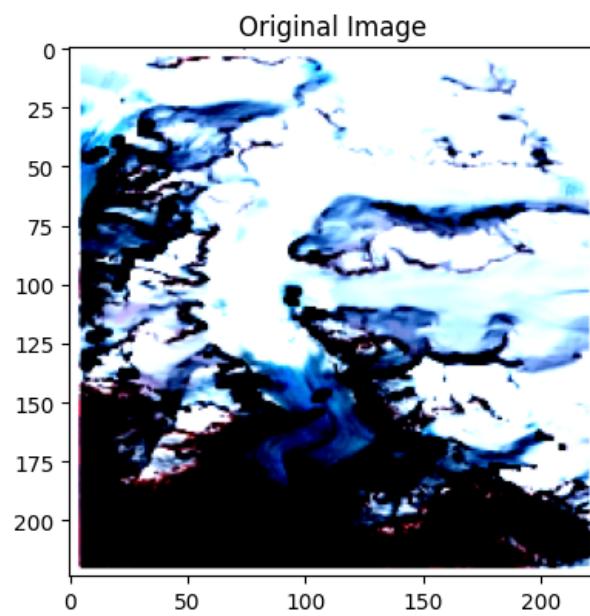
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type(bool) for binary data



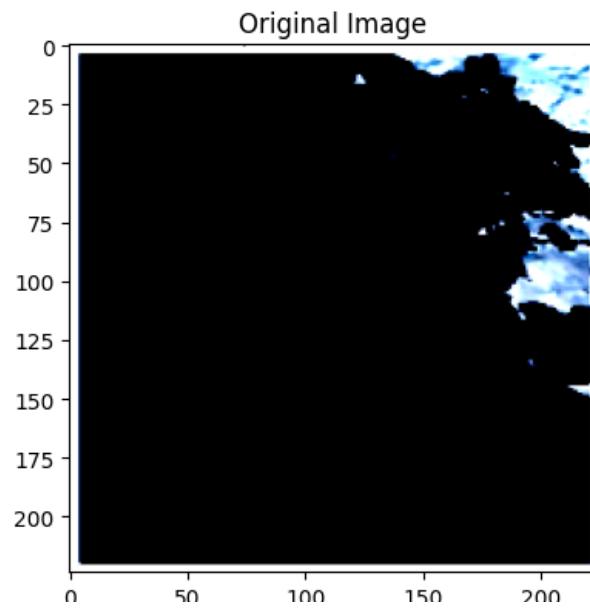
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type(bool) for binary data



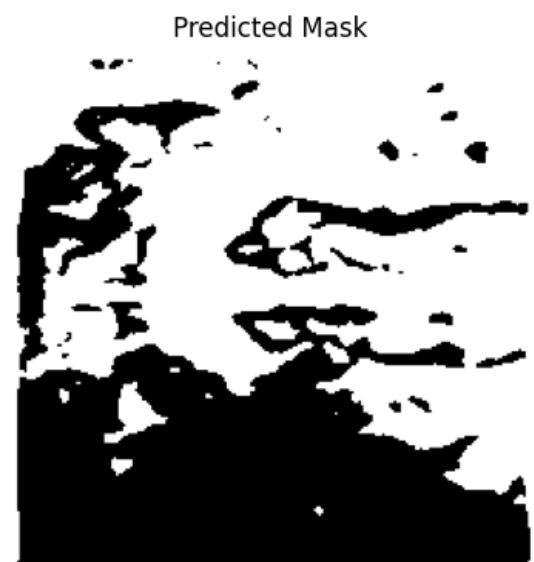
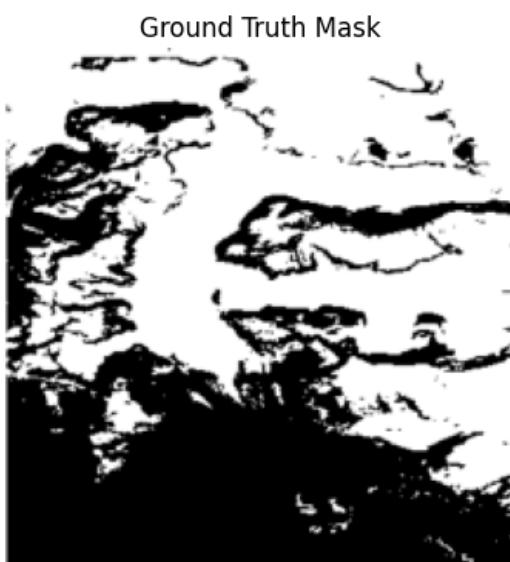
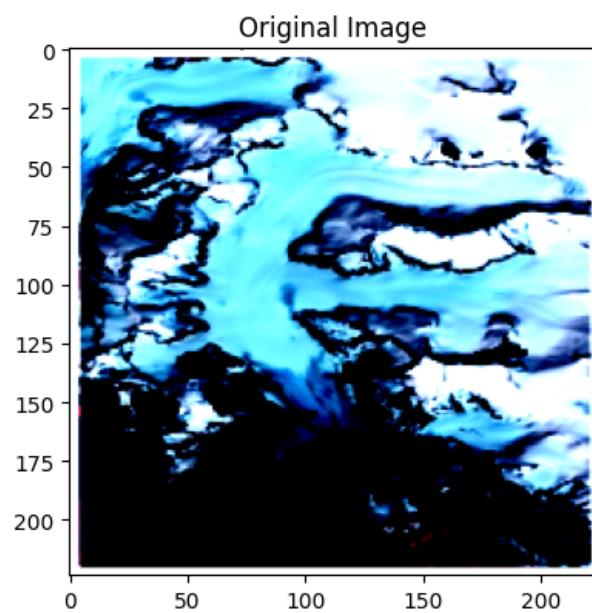
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type(bool) for binary data



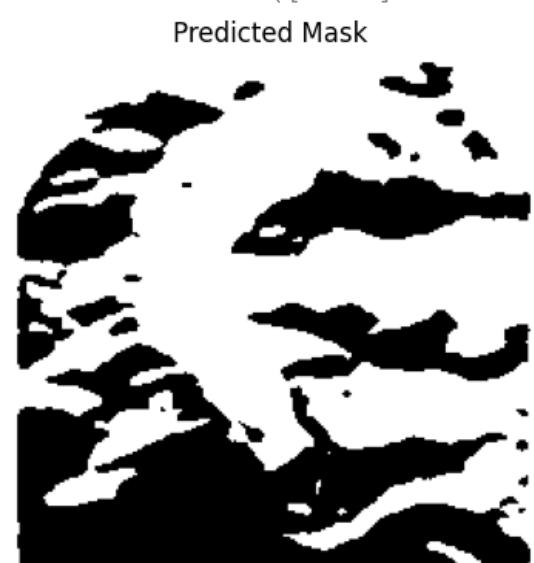
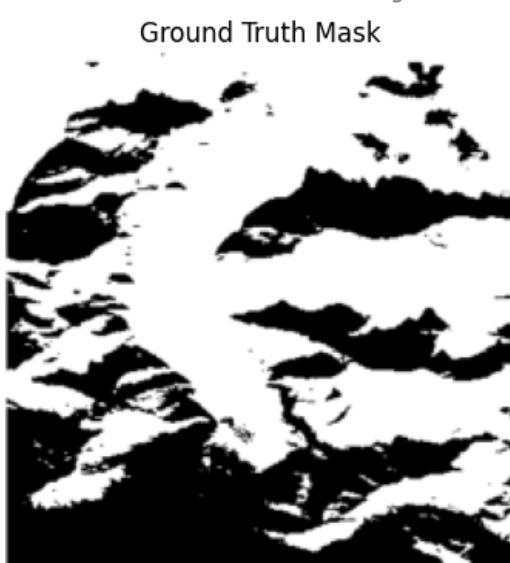
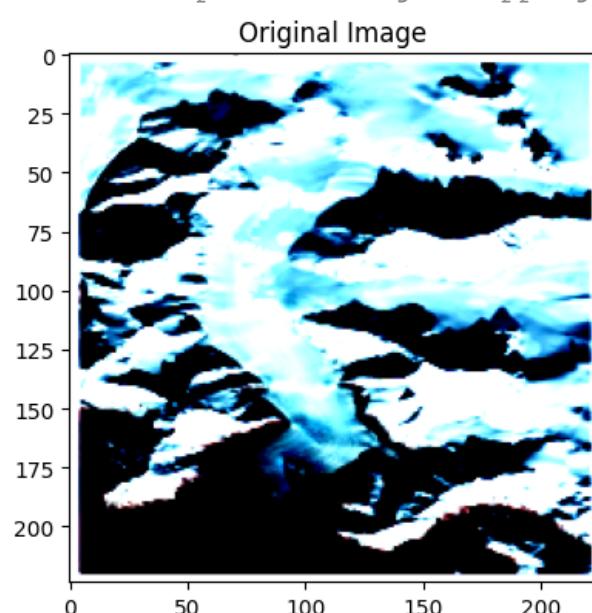
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) and type(bool) for binary data



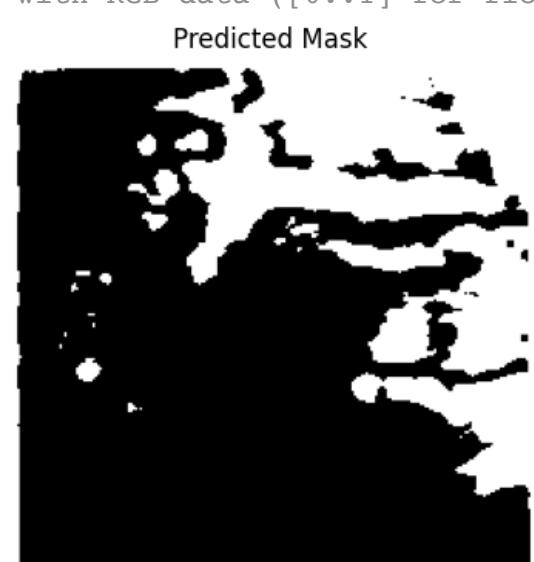
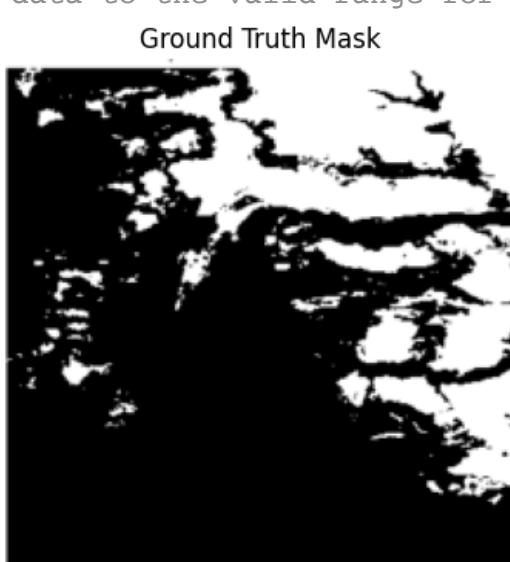
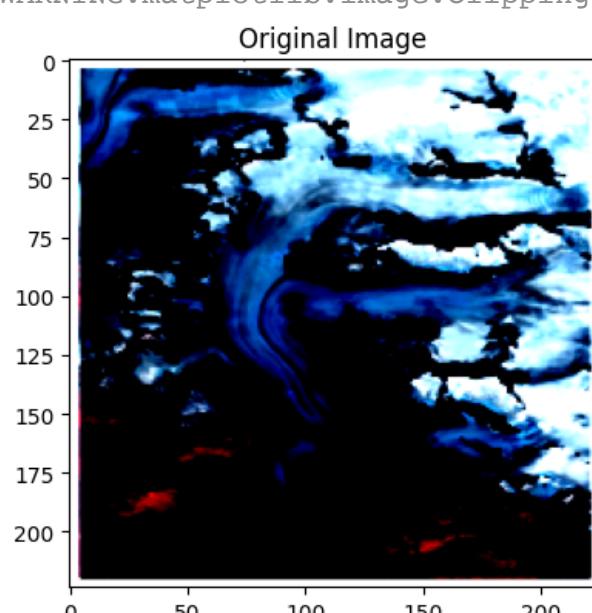
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar.



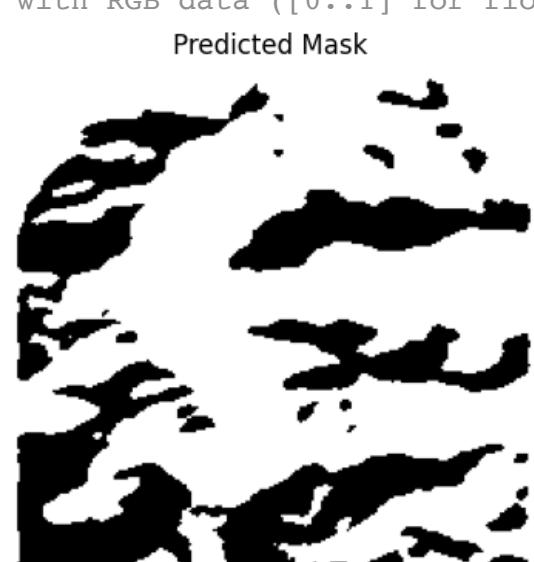
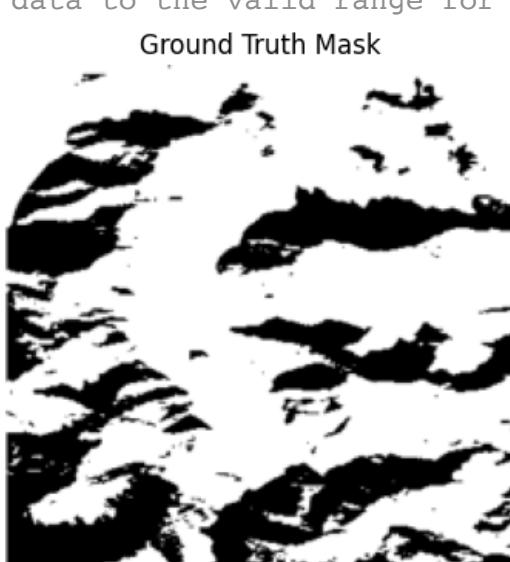
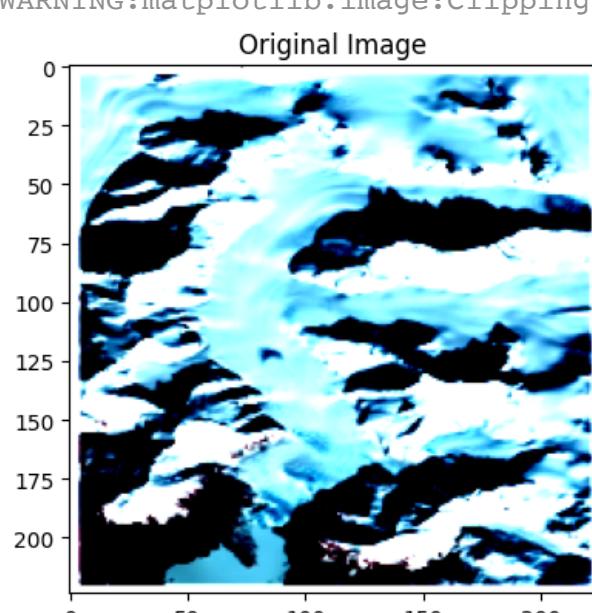
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar.



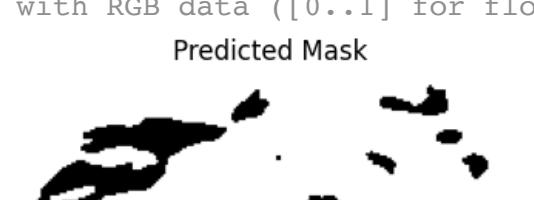
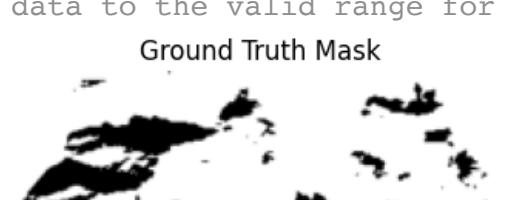
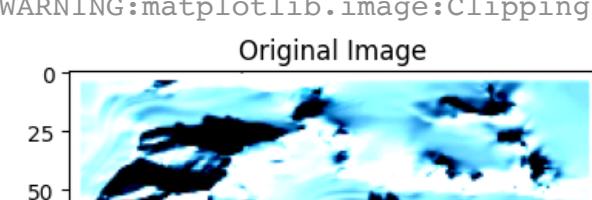
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar.

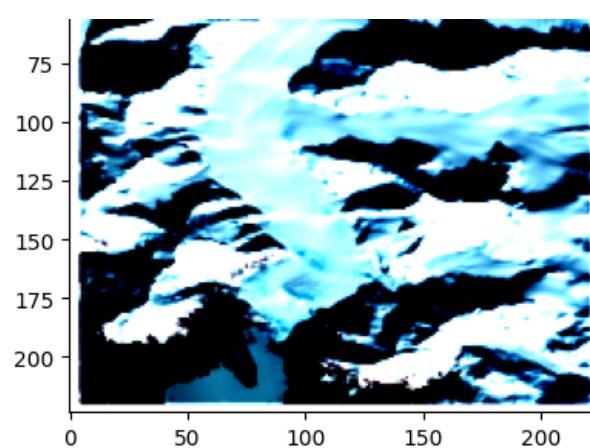


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar.

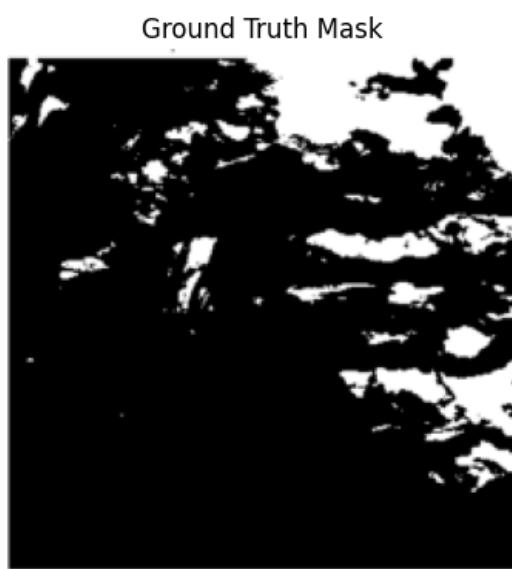
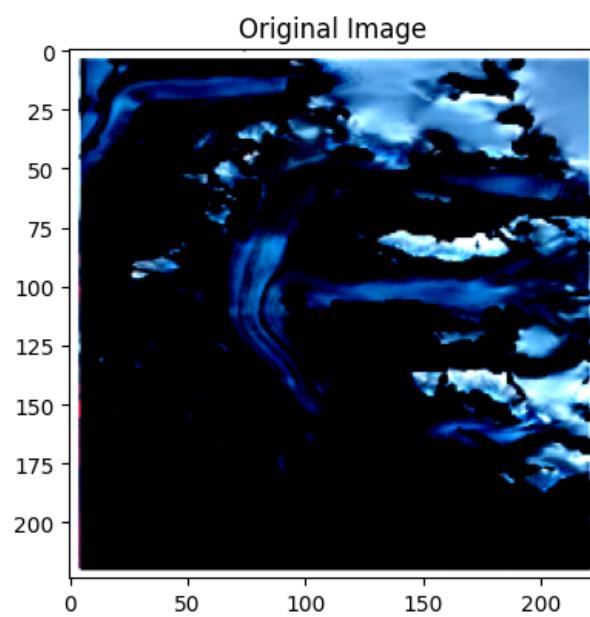


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for uint8 and similar.

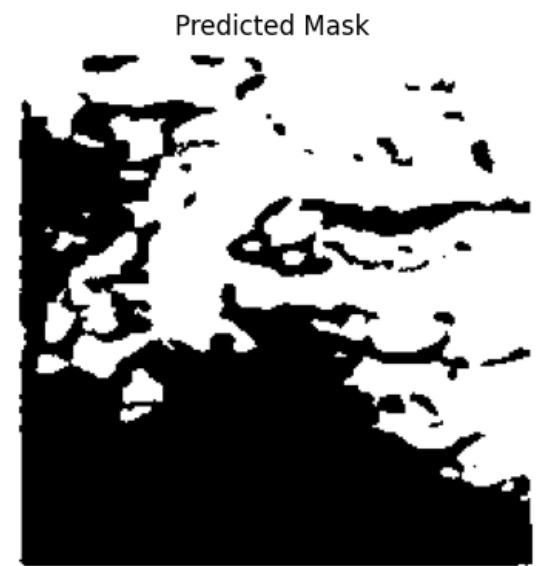
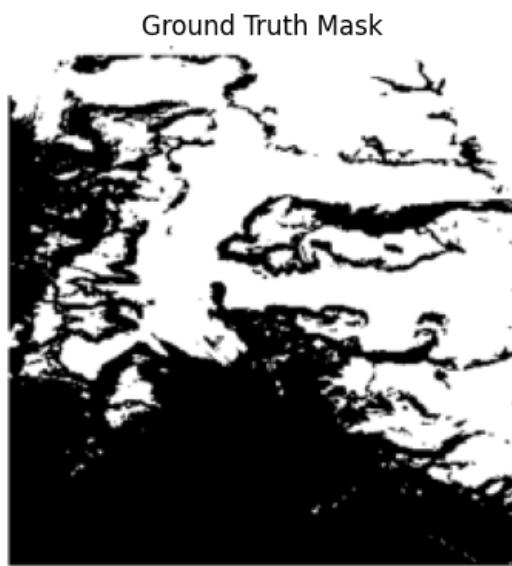
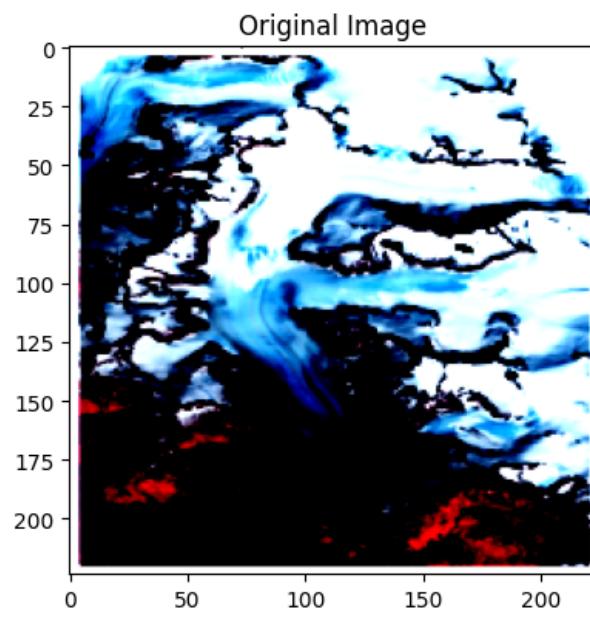




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float RGB data)



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float RGB data)

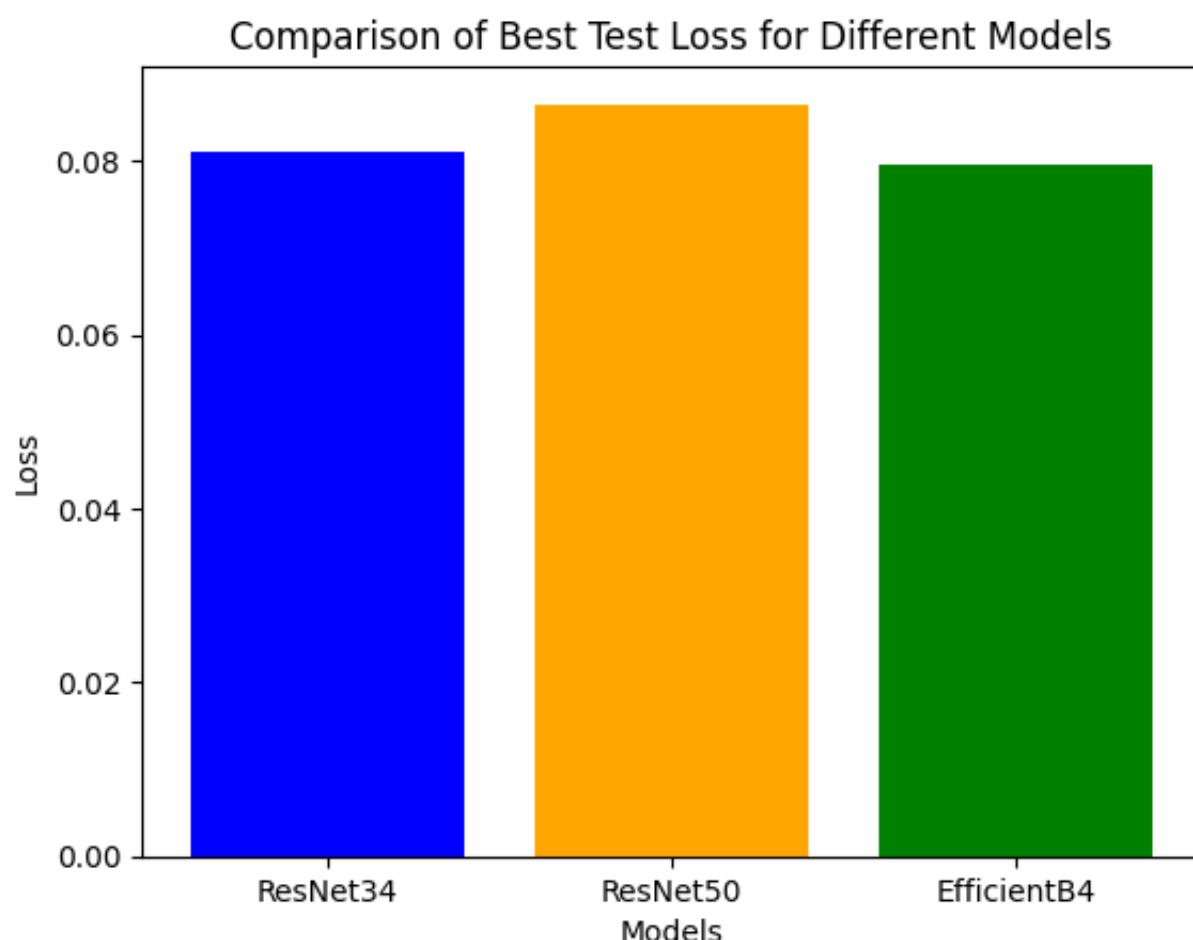


## Comparison of All Three Models | Best Performing Model based on Loss Scores - ResNet34

```

1 # # Performance explanation using BarPlot:
2 # - ResNet34 performed best as it has sufficient capacity for the dataset's complexity while being
3 # - EfficientNet-B4 performed well due to its advanced design but may require more data or hyperpar
4 # - ResNet50 underperformed because its higher complexity likely led to overfitting on the available
5 loss_values = [resnet34_best_test_loss, resnet50_best_test_loss, efficientb4_best_test_loss]
6 models = ['ResNet34', 'ResNet50', 'EfficientB4']
7 plt.bar(models, loss_values, color=['blue', 'orange', 'green'])
8 plt.title('Comparison of Best Test Loss for Different Models')
9 plt.xlabel('Models')
10 plt.ylabel('Loss')
11 plt.show()

```



```

1 # printing the individual values for clear comparison, as explained in the previous cell, resnet34 :
2 resnet34_best_test_loss, resnet50_best_test_loss, efficientb4_best_test_loss
(0.08100565274556477, 0.08649138609568278, 0.07955124974250793)

```

```

1 # comparing the predicted output of the best weights saved for all the three models on one sample to
2 # On closer observation, it can be viewed that ResNet34 segments minor regions better followed by ef
3 test_dataset = SegmentationDataset([image_paths[81]], [mask_paths[81]])
4 test_loader = DataLoader(test_dataset, batch_size=1, shuffle=True)
5
6 for encoder_name in ["resnet34", "resnet50", "efficientnet-b4"]:
7     print(f'For encoder - {encoder_name}')
8     model = smp.Unet(
9         encoder_name=encoder_name,
10        encoder_weights="imagenet",
11        in_channels=3,
12        classes=1
13    )
14    model = model.to(device)
15    model.load_state_dict(torch.load(f'Best_model_{encoder_name}.pth'))
16    test_image_plot(test_loader, model)

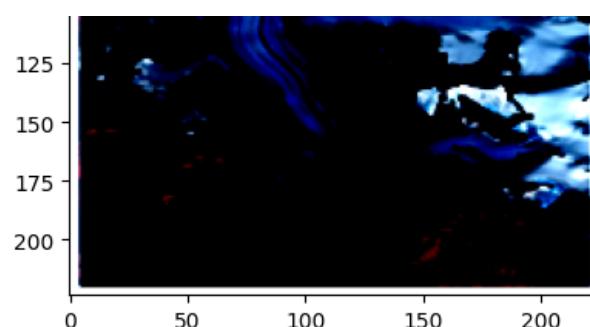
```

```

For encoder - resnet34
<ipython-input-27-62e9b54a6186>:15: FutureWarning: You are using `torch.load` with `weights_only=False` (
  model.load_state_dict(torch.load(f'Best_model_{encoder_name}.pth')))
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float

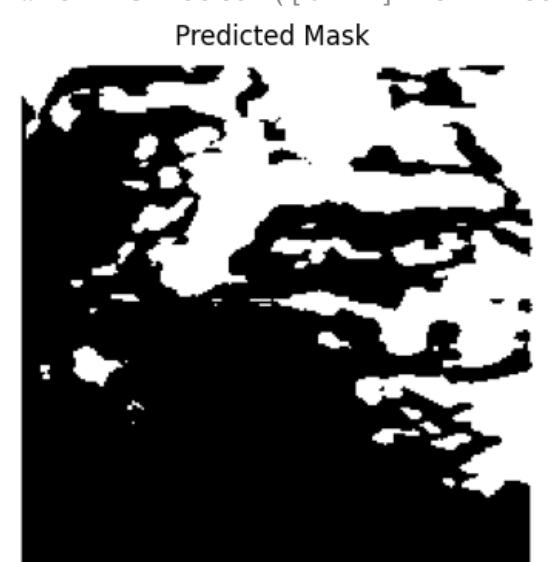
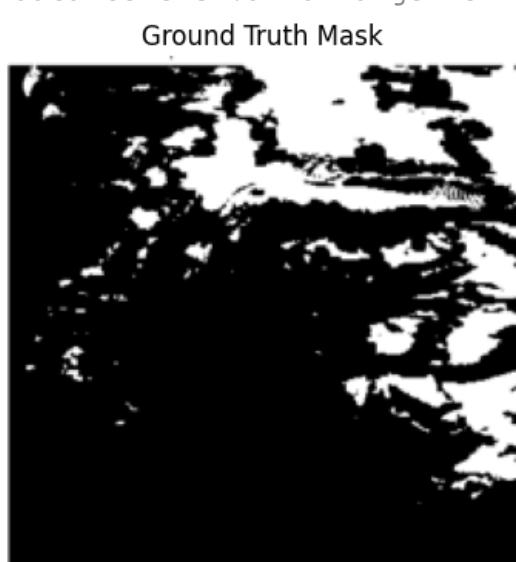
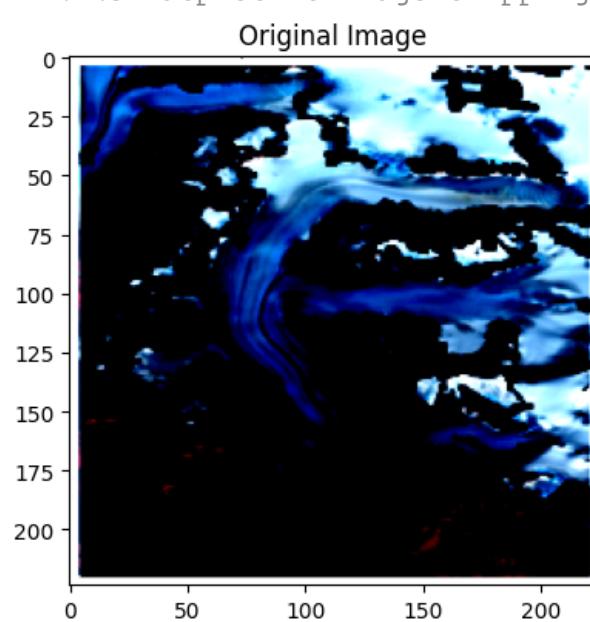
```





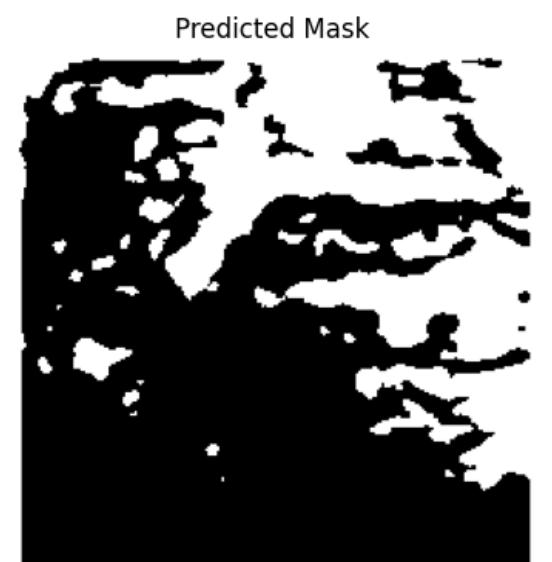
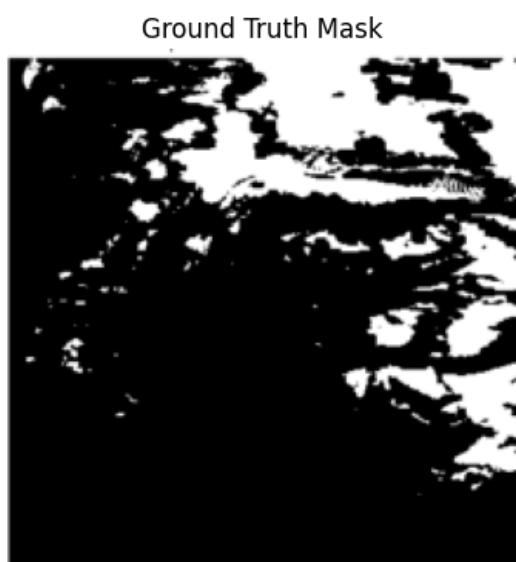
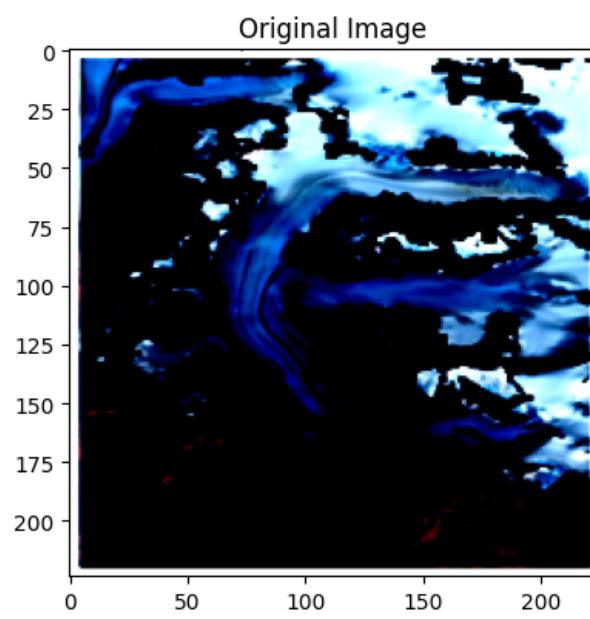
For encoder - resnet50

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



For encoder - efficientnet-b4

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for float32) or [0..255] for integer types



## Deployment Considerations

```
1 Here, the first step for deploying any model would be exporting it as .pth file, which is storing i+
2 which as seen in the above code, I have already saved the models.
3
4 I have created a global loss variable, which keeps track of the lowest loss achieved across all cro+
5 saved the best one, and everytime new minimum loss would be obtained, model for that step using tor+
6 will be saved.
7
8 Then there are two types of deployment possible:
9
10 (a) Cloud-based Deployment:
11
12 Possible Cloud Platforms: AWS(Sage Maker), Azure, and Google Cloud
13 I prefer Google Cloud because of its integration with BigQuery for data analytics, and support for (
14 the primary source of my dataset.
15
16 Steps to Deploy on Google Cloud Platform:
17
18 1) Set up your project and environment.
19 2) Create an image classification dataset, and import images.
20 3) Train an AutoML image classification model.
21 4) Evaluate and analyze model performance.
22 5) Deploy a model to an endpoint, and send a prediction.
23 6) Clean up your project.
24
25 I got these steps from this tutorial: https://cloud.google.com/vertex-ai/docs/tutorials/image-classification
26
27 Alternatively, we can also choose to deploy it as a Docker:
28
29 A sample docker file would look something like this:
30
31 FROM pytorch/pytorch:latest
32
33 WORKDIR /app
34
35 COPY . /app
36
37 RUN pip install -r requirements.txt
38
39 CMD ["python", "inference_server.py"]
40
41 (b) Deploy as an API:
42
43 We may use Flask/Django to create an UI, and connect the backend with either our cloud platform or o+
44 The reason I am suggesting Flask/Django is because both are Python-based and hence easy to integrate.
45 For large-scale deployment, I would suggest Django because of its additional security constraints, whereas F
46 Flask would be the easy-to-go option.
47
48 (c) Deploy on Edge-Device:
49
50 We may deploy the model on edge-devices like Raspberry-Pi or Jetson, though I would suggest using Je+
51 because previously I have worked with both and from my experience I can say that raspberry-pi's die
52 neural networks, while Jetsons provide performance similar to small GPUs.
```

## ✓ Steps to make my model real-time:

```
1 1) Create data-fetching code for automatically retrieving latest images from GoogleEarthEngine for a  
2 Here I have manually given the time-frame but instead of that I can give a specific range and retrieve  
3 Though some additional problems may arise: it is not necessary that everytime a file would arrive or  
4 files for specific time periods, so I will have to sort all the files in the system as per datetime  
5 keeping a fixed number.  
6 2) Run that data-fetching code on cron-job, it is a linux tool, which runs a described pipeline even  
7 3) Then run the trained model in eval mode to generate segmented masks.  
8 4) Present the output to a user using some GUI, or store the segmented masks on some permanent storage  
9  
10  
11  
12 One more additional feature can be added:  
13 1) Few features of the images may change over-time, hence our model need to capture those  
14 new characteristics, and for that we should incorporate continual learning/incremental learning mechanism  
15 in our model.  
16 Challenges in Continual Learning:  
17 Catastrophic Forgetting: The phenomenon where a model forgets previously learned knowledge when trained  
18 Scalability: Handling an ever-increasing amount of data without requiring large storage or computation  
19 Data Distribution Shift: Adapting to changes in the underlying data distribution over time (non-stationary)  
20 Efficient Memory Usage: Ensuring the model doesn't require storing all previous data, which might be  
21 Methods to Incorporate Continual Learning:  
22 Regularization-based methods (EWC, MAS).  
23 Replay-based methods (Experience Replay, Generative Replay).  
24 Dynamic architecture methods (Progressive Networks, PathNet).  
25 Meta-learning methods (MAML, Reptile).  
26 Neuro-inspired methods (Self-Organizing Maps, Hebbian Learning).  
27 Knowledge distillation methods for transferring knowledge across tasks
```

## ▼ Ethical Consideration

## 1 # Fairness:

2 Bias in Training Data: Here the model is trained on just the images from the data obtained over the areas. If the training data is biased towards specific regions or types of glaciers, the model may fail to generalize well to other areas. This is particularly problematic if the model is applied to less-represented areas. Hence we need to make sure that the training datasets are diverse, representing a wide range of outcomes. We can use techniques like adversarial training or fairness constraints to identify any biases that may affect the model's performance.

7

## 8 # Societal Impact:

9

10 The use of machine learning models in glacier segmentation has significant societal implications, particularly in climate change and environmental monitoring. As glaciers play a critical role in the Earth's water cycle, accurate modeling of glacier dynamics can help predict future water availability, potential natural disasters, and even climate change. While the societal benefits are clear, several ethical issues need to be considered.

14

## 15 # Accuracy vs. Accountability:

16 If the model is used to make important environmental or policy decisions (e.g., regarding glacier management or climate change prevention), errors in the segmentation model could have large-scale consequences. The ethical dilemma lies in balancing accuracy and transparency while avoiding overreliance on the AI system for high-stakes decision-making. Hence it is important to implement continuous model monitoring, and human-in-the-loop mechanisms to ensure that the model's predictions are accurate and transparent.

## 21 # StakeHolder Effects:

22 Various stakeholders are involved in or affected by the development and deployment of glacier segmentation models, including government organizations, policy makers, local communities, and industries (e.g., agriculture, tourism, or hydroelectric power generation).

25 The information derived from glacier segmentation models may be used for resource management or environmental monitoring. If this information is shared responsibly, it could lead to better management of local resources and environmental degradation. However, if local communities are not included in the decision-making process, and any changes informed by model results are not communicated effectively, it could lead to resistance and non-compliance.

29 Machine learning systems for environmental monitoring are often developed and deployed by large research institutions or commercial entities, which may not always prioritize the interest of marginalized communities. These groups might lack access to the technology or the benefit of collaborative partnerships to ensure that smaller organizations, local governments, and vulnerable communities are included in the decision-making process.

32

## 33 # Transparency and Accountability:

34 Machine learning models, especially deep learning models, are often seen as "black boxes" because it is difficult to understand how they arrive at their conclusions. If a glacier segmentation model misidentifies certain areas as glaciers or non-glaciers, it is important to improve transparency, allowing stakeholders to understand how the model works and how decisions are made.

37

38 To implement this, softwares like SHAP/LIME can be used.

39

## 40 #Data Ownership and Consent:

41 In cases where drone-based or sensor-based data is used to capture information on specific glaciers, it is important to understand who owns the data and how it is shared.

43 Without proper consent, data could be used in ways that violate privacy rights or lead to the exploitation of vulnerable communities. For example, if a glacier segmentation model误识别了某些区域为冰川，而这些区域实际上不是冰川，那么该模型可能会导致资源的浪费或错误的决策。

44 Here, for this project all the rights are reserved with GoogleEarthEngine.

1 Start coding or generate with AI.

